

Objective: To create, train and test a Convolutional Neural Network for the task of semantic segmentation using FCN32/FCN16 architecture on Kitti Semantic Segmentation dataset.

1. Introduction:

Convolution Neural Networks (CNN) derive its strength from learning the filters (Kernels) at each layer to extract features from the image and updating the weights iteratively using backpropagation. CNNs were extensively used for image classification tasks where the target is class labels. But for tasks like Image Segmentation where the target is a structured output where every pixel in the image is labeled, the conventional convolution networks had to be modified. Inorder to develop an end-to-end network for Semantic Segmentation, Fully Convolutional Network (FCN) was built in 2015 which achieved better results on PASCAL VOC 2012 dataset. Figure 1 shows the network architecture of FCN32.

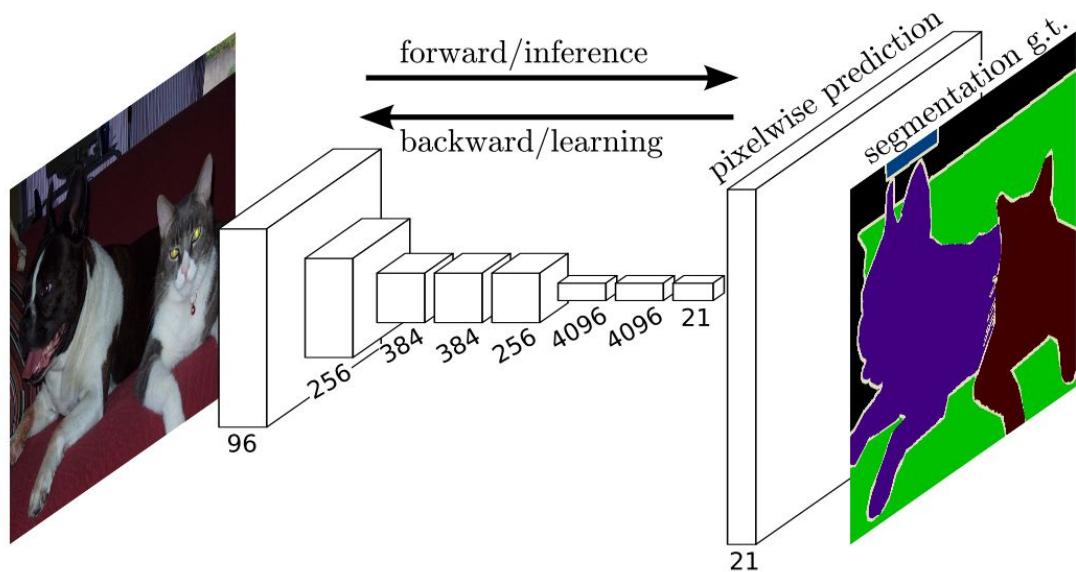


Figure 1: FCN Architecture

I. Classifier to dense FCN:

FCN adopts transfer learning by using pre-trained networks for image classification tasks like VGG-16, AlexNet, GoogLeNet. In this assignment, we are using the pre-trained model of VGG-16 network. We discard the last classifier layer and replace them with convolutional layers. In the last layer, a 1×1 convolution layer with 35 channels to predict the scores of each class at each output location followed by a deconvolution layer to upsample the coarse outputs to pixel dense output.

II. FCN-16 to obtain finer predictions:

By changing the architecture above the pool-5 layer by replacing the existing filters with smaller filters, we get finer predictions. FCN-16 is realized by combining the predictions from both the final layer and the pool4 layer, at stride-16. Figure 2 shows the modifications to original network.

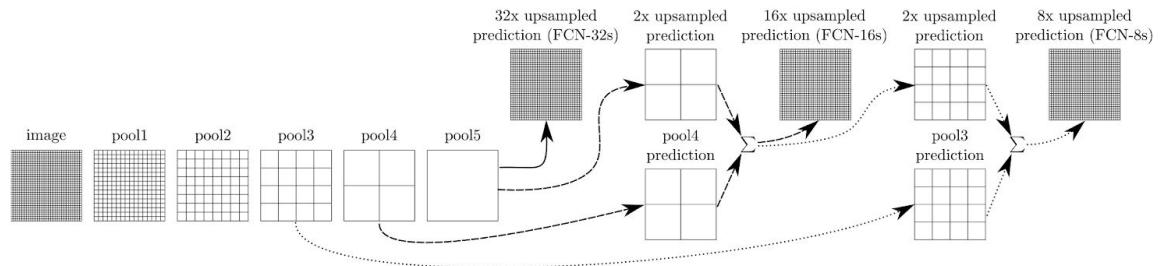


Figure 2: Modified FCN architecture - FCN-16 and FCN-8

2. Summary of the Program:

2.1 Training Information: A brief description of the training process and the parameters are listed below,

- a) **Network architecture:** We use pre-trained VGG-16 architecture and convert the last fully connected layers to convolutional layers to get pixel-wise predictions.
- b) **Cost function:** Cross entropy loss is used during the training process on training and validation set.
- c) **Evaluation metric:** We use pixel level IoU on all the testset images and mIoU (mean IoU) on all testset images to measure the generalization ability of the network.
- d) **Preprocessing:** All the input images are resized to 256*516 using BICUBIC interpolation technique and all labels are resized to the same dimension using NEAREST interpolation technique.
- e) **Training parameters:** We used SGD and Adam optimizer for training with a learning rate of 0.001, batch size 5, momentum 0.99 and the model is trained for 30-40 epochs.

2.2 Program and dataset information: The code for data loader, training and evaluation is available in the notebook *FCN_semantic_segmentation.ipynb*. The dataset is available in the *data_semantics* folder where we only used 200 training samples and split the data into training, validation and test sets. All the trained models are also included in the zip file *Models*.

2.3 Source Listing:

A. Data Loader code:

```
class semantic_dataset(Dataset):
    def __init__(self, split = train, transform = None):
        self.split = split
        self.transform = transform
        self.img_list = image_train
        self.mask_list = label_train

    def __len__(self):
        return(len(self.img_list))

    def __getitem__(self, idx):
        img = cv2.imread(self.img_list[idx])
        img = cv2.resize(img, (1242, 376))
        mask = None
        if self.split == 'train':
            mask = cv2.imread(self.mask_list[idx],
cv2.IMREAD_GRAYSCALE)
            mask = cv2.resize(mask, (512, 256),
interpolation=cv2.INTER_NEAREST)
            #mask = self.encode_segmap(mask)
            assert(mask.shape == (256, 512))

        if self.transform:
            img = self.transform(img)
            assert(img.shape == (3, 256, 512))
        else :
            assert(img.shape == (256, 512, 3))

        if self.split == 'train':
            return img, mask
        else :
            return img
```

B. Preprocessing and loading data:

```

transformation = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize((256,512),interpolation=PIL.Image.BICUBIC),
    transforms.ToTensor(),
    transforms.Normalize(mean      = [0.35675976,      0.37380189,
0.3764753], std = [0.32064945, 0.32098866, 0.32325324])
])
trainset = semantic_dataset_train(split = 'train', transform =
transformation)
validationset = semantic_dataset_validation(split = 'train',
transform = transformation)
testset = semantic_dataset_test(split = 'train', transform =
transformation)

trainImages = DataLoader(trainset, batch_size = 5, shuffle =
True)
valImages = DataLoader(validationset, batch_size = 5, shuffle =
=True)
testImages = DataLoader(testset, batch_size = 1, shuffle =
True)

```

C. FC32 Network:

```

#FC32 architecture
vgg16_mod = vgg16(pretrained=True)
for param in vgg16_mod.features.parameters():
    param.requires_grad = False
class FCN32(nn.Module):
    def __init__(self):
        super(FCN32, self).__init__()
        self.features = vgg16_mod.features
        self.classifier = nn.Sequential(
            nn.Conv2d(512, 4096, 7),
            nn.ReLU(inplace=True),
            nn.Dropout2d(),
            nn.Conv2d(4096, 4096, 1),
            nn.ReLU(inplace=True),

```

```
        nn.Dropout2d(),
        nn.Conv2d(4096, 35, 1),
        nn.ConvTranspose2d(35, 35, 224, stride=32)
    )
    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x
fcn = FCN32()
print(fcn)
input = torch.randn(5,3,256,512)
output = fcn(input)
print(output.shape)
```

D. FC-16 architecture

```
#FC16 architecture
vgg16_mod = vgg16(pretrained=True)
for param in vgg16_mod.features.parameters():
    param.requires_grad = False
class FCN16(nn.Module):
    def __init__(self):
        super(FCN16, self).__init__()
        self.features = vgg16_mod.features
        self.classifier = nn.Sequential(
            nn.Conv2d(512, 4096, 7),
            nn.ReLU(inplace=True),
            nn.Conv2d(4096, 4096, 1),
            nn.ReLU(inplace=True),
            nn.Conv2d(4096, 35, 1)
        )
        self.score_pool4 = nn.Conv2d(512, 35, 1)
        self.upscore2 = nn.ConvTranspose2d(35, 35, 14, stride=2,
bias=False)
        self.upscore16 = nn.ConvTranspose2d(35, 35, 16, stride=16,
bias=False)
```

```
def forward(self, x):
    pool4 = self.features[:-7](x)
    pool5 = self.features[-7:](pool4)
    pool5_upscored = self.upscore2(self.classifier(pool5))
    pool4_scored = self.score_pool4(pool4)
    combined = pool4_scored + pool5_upscored
    res = self.upscore16(combined)
    return res

fcn16 = FCN16()
print(fcn16)
input = torch.randn(5,3,256,512)
output = fcn16(input)
print(output.shape)
```

E. Training:

```
def training(fcn, criterion, optimizer, num_epochs):

    val_loss = []
    train_loss = []
    for epoch in range(num_epochs):
        fcn.train()
        val_running_loss = 0
        train_running_loss = 0
        fcn16.train()
        for i, trainData in enumerate(trainImages, 0):
            imgs, labels = trainData
            labels = labels.type(torch.LongTensor)
            #set the gradient to zero
            optimizer.zero_grad()
            outputs = fcn(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            train_running_loss += loss.item()
            optimizer.step()

        train_loss.append(train_running_loss)
```

```
fcn16.eval()
with torch.no_grad():
    for k, dat1 in enumerate(valImages):
        # Get image, label pair
        inputs1, labels1 = dat1
        labels1 = labels1.type(torch.LongTensor)
        # Predicting segmentation for val inputs
        outputs1 = fcn(inputs1)

        # Compute CE loss and aggregate it
        loss1 = criterion(outputs1, labels1)
        val_running_loss += loss1.item()

    val_loss.append(val_running_loss)
    print("Epochs: ", epoch, "Training Loss: ", train_running_loss, "val_running_loss: ", val_running_loss)
return train_loss, val_loss

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(fcn16.parameters(), lr=0.001,
betas=(0.9, 0.999), eps=1e-08, weight_decay=0.0001,
amsgrad=False)
train_loss, val_loss = training(fcn16, criterion, optimizer,
num_epochs=40)
torch.save(fcn16, 'drive/My
Drive/CSCI-677-HW5/data_semantics/ModelFCN16_Adam_dropout.pt')
```

F. Calculating Mean IoU and Dice coefficient (F1-score)

(This is calculated using a *confusion matrix*. Pixel level IoU for each class is used. Complete code is attached at the end.)

```
#Calculating IoU and mean IoU
def get_mean_iou(conf_mat, multiplier=1.0):
    cm = conf_mat.copy()
    np.fill_diagonal(cm, np.diag(cm) * multiplier)
    inter = np.diag(cm)
    gt_set = cm.sum(axis=1)
    pred_set = cm.sum(axis=0)
    union_set = gt_set + pred_set - inter
    iou = inter.astype(float) / union_set
    mean_iou = np.nanmean(iou)
    return mean_iou
overall_conf_mat = np.zeros((35, 35))
with torch.no_grad():
    for k, dat1 in enumerate(testImages):
        # Get image, label pair
        inputs1, labels1 = dat1
        # Predicting segmentation for val inputs
        outputs1 = fcnn16(inputs1)

        preds = torch.argmax(outputs1, dim=1).detach().numpy()
        gt = labels1.detach().numpy()

        # Compute confusion matrix
        conf_mat = confusion_matrix(y_pred=preds.flatten(),
                                     y_true=gt.flatten(), labels=list(range(35)))
        overall_conf_mat += conf_mat

    print('DICE score: {}'.format(np.round(get_mean_iou(conf_mat=overall_conf_mat,
                                                       multiplier=2.0), 2)))
    print('Mean IOU score: {}'.format(np.round(get_mean_iou(conf_mat=overall_conf_mat),
                                              2)))
```

3. Experimental Results:

(Results on all 30 test images along with the code is attached in the end)

FCN-32 is trained with two variations, one with SGD optimizer and the second with Adam optimizer, FCN-16 is trained with Adam optimizer and the performance of each model is evaluated and discussed later.

A. FCN-32 with Adam optimizer:

Result:

Dice score (F1-score): 0.45

Mean IoU score: 0.37

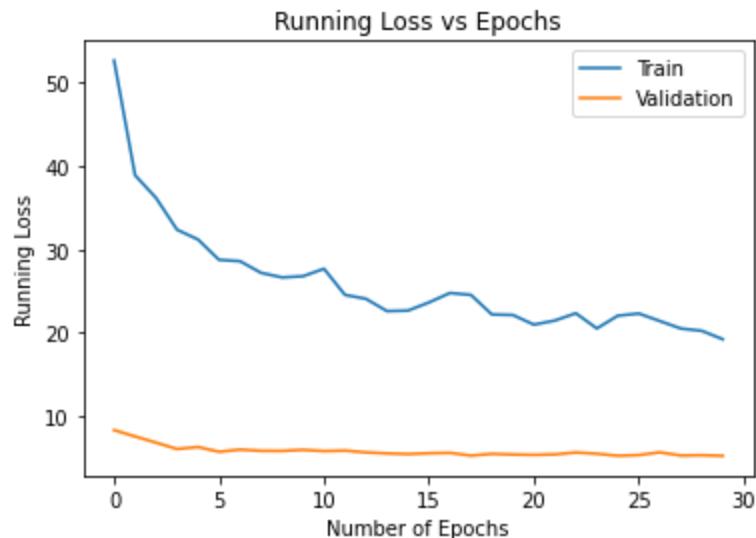
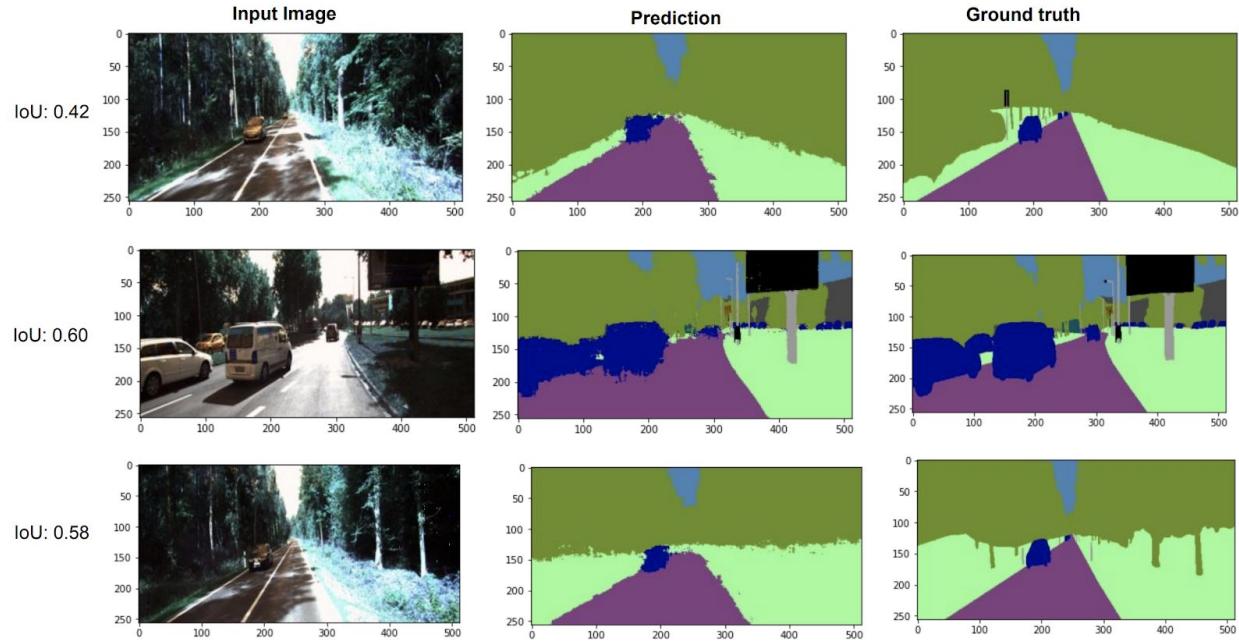


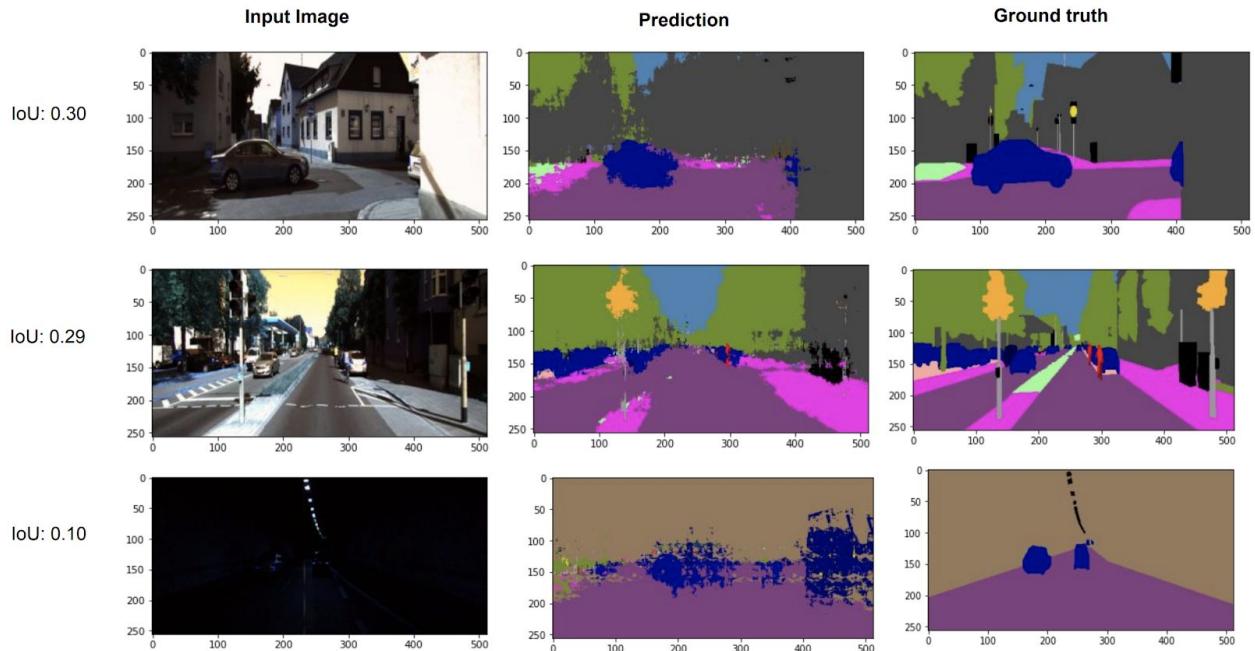
Figure 3: Evolution of training and validation loss on each epoch.

Visualization of results:

(Top 3 successful results: Based on pixel level IoU using confusion matrix)



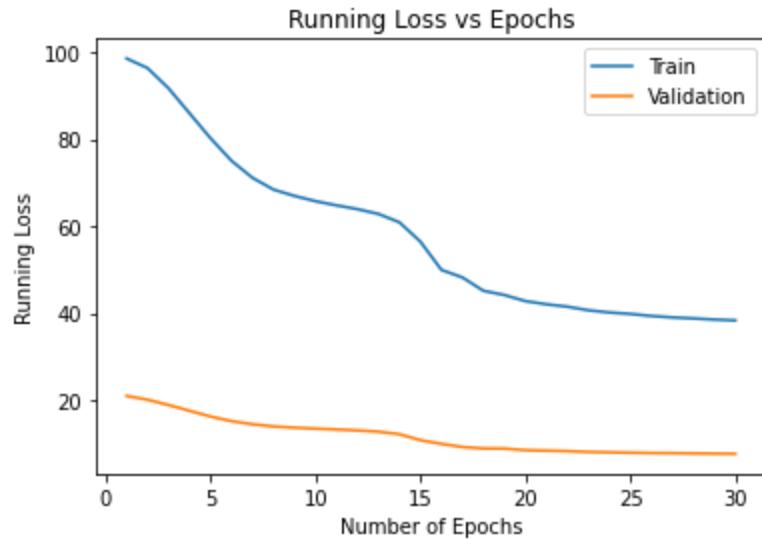
(Top 3 failed results: Based on pixel level IoU using confusion matrix)



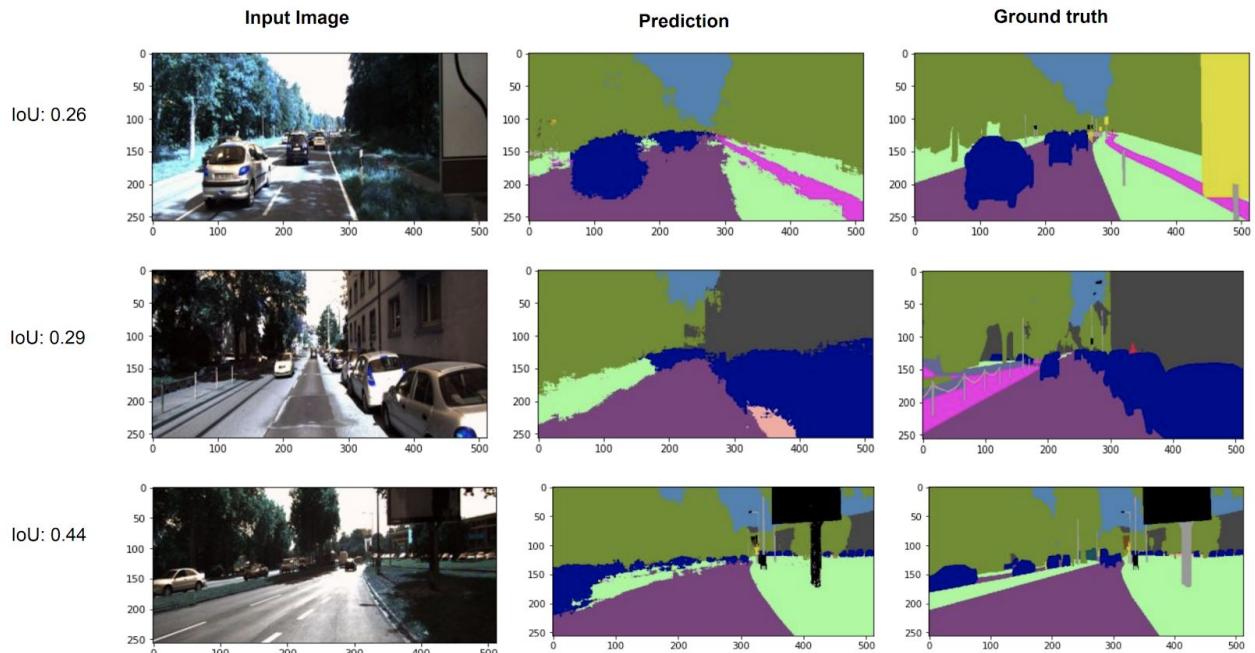
B. FCN-32 with SGD optimizer:**Result:**

Dice score (F1-score): 0.35

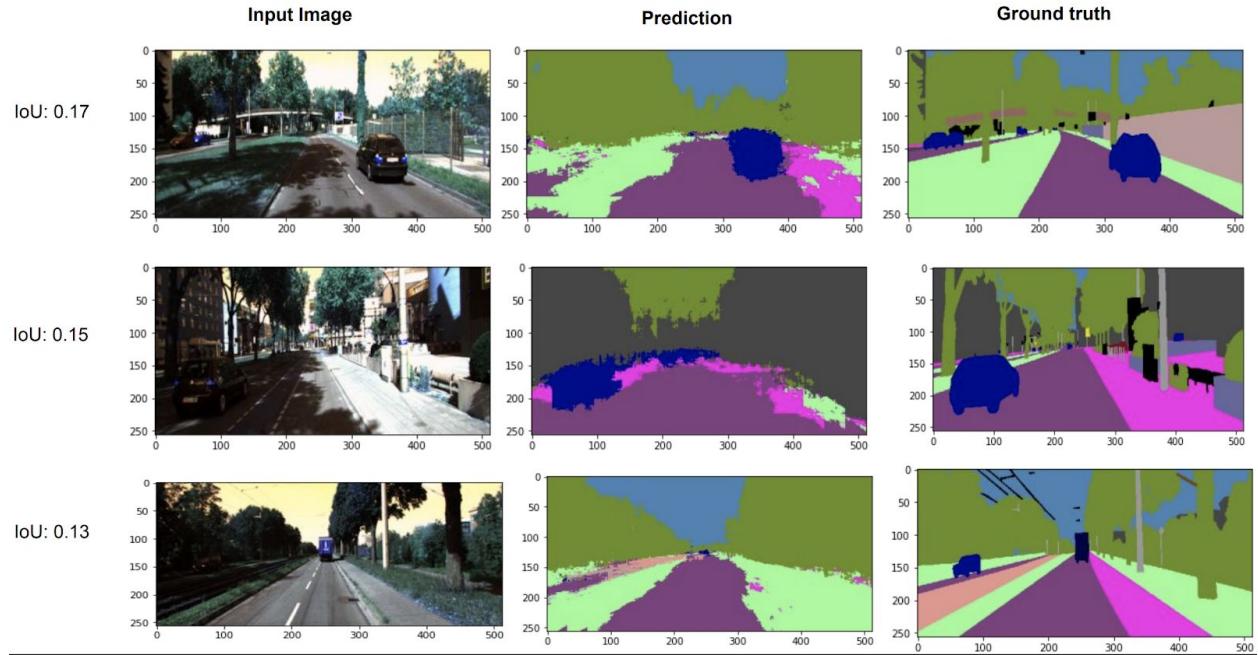
Mean IoU score: 0.23

**Figure 4: Evolution of training and validation loss on each epoch.**

(Successful results with pixel level IoU using confusion matrix)



(Failed results with pixel level IoU using confusion matrix)



C. FCN-16 with Adam optimizer:

Result:

Dice score (F1-score): 0.58

Mean IoU score: 0.43

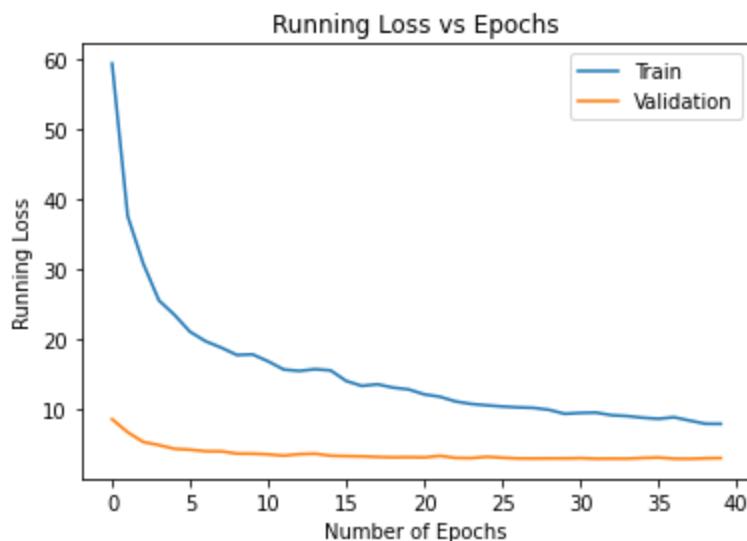
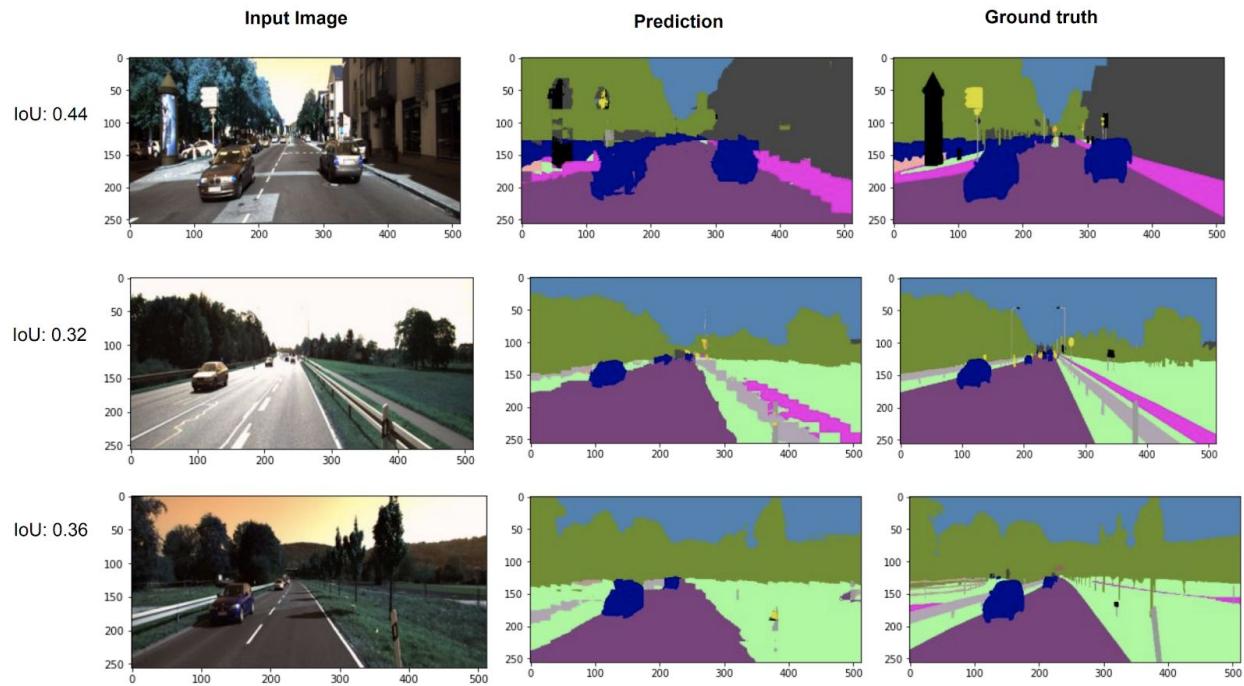
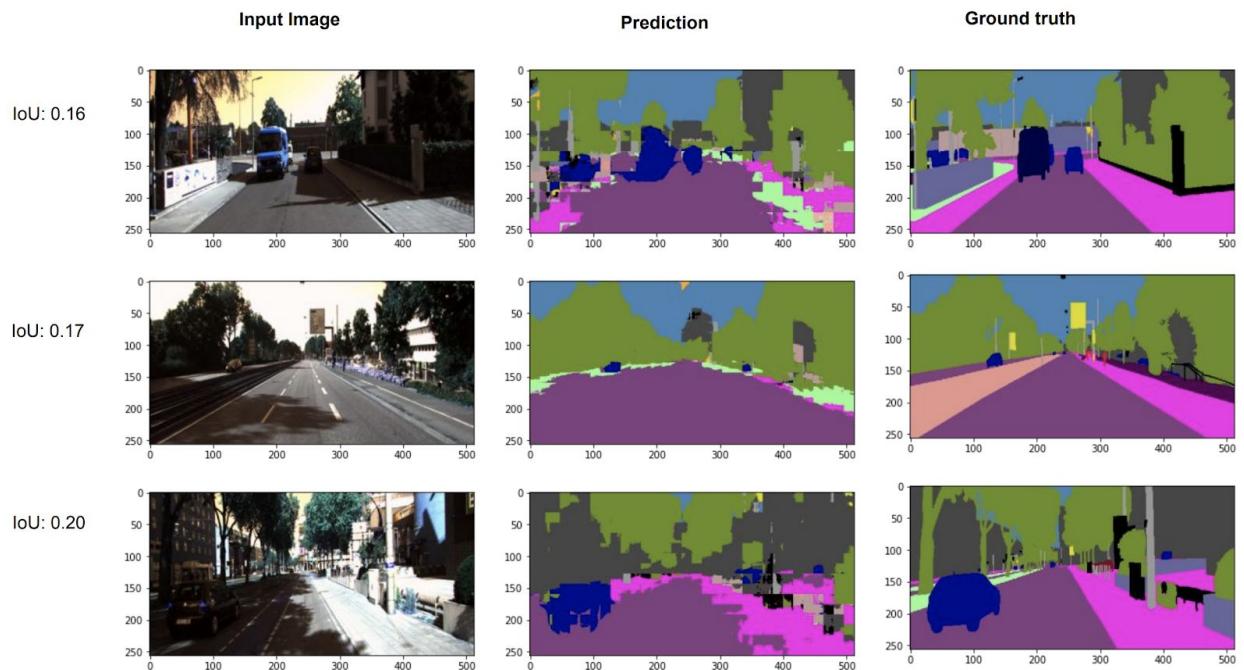


Figure 5: Evolution of training and validation loss on each epoch.

(Successful results with pixel level IoU using confusion matrix: FCN-16)



(Failed results with pixel level IoU using confusion matrix:FCN-16)



4. Discussion on Results:

FCN-32 is trained using Adam and SGD optimizers with a learning rate of 0.001 for 30 epochs. From the above results, it is evident that the results of Adam optimizer is better than SGD optimizer. Observing at the loss curve, we can say that the model trained using SGD optimizer seems to underfit the data. This argument is supported by the low mean-IoU score of 0.23. However, we obtained a better performance for the model trained using Adam optimizer and we got moderate improvement in mean-IoU score (0.37).

FCN-16 is trained using Adam optimizer for 40 epochs. Based on the loss curve, we can say that the model did not overfit the data. We tested the model on the test set and we obtained a best mean-IoU score of 0.43. Since FCN-16 uses upsampled output from the pool-4 layer combined with output from pool-5 layer, we get finer predictions.

Results of semantic segmentation along with the ground truth is shown in the above figures for each model. Few successful results based on the m-IoU score are displayed and to evaluate the shortcomings of the model, few failed examples are also shown. The model is performing extremely well on images which have proper illumination with little or no occlusion. However, the model terribly failed for the images which have shadows in it and have significant occlusions. Also, few examples are shown where the far objects are missed in the predictions.

FCN-16 produced better results by giving finer predictions compared to FCN-32. This can be seen in the visual results where classes like vehicles, vegetation, road, sky were predicted accurately in both FCN-32 and FCN-16. However, classes like traffic lights, sidewalk, and poles are missed in FCN-32 and are predicted correctly in FCN-16.

5. Appendix:

(Complete code with results)

```
from google.colab import drive
drive.mount('/content/drive',force_remount=True)
```

Mounted at /content/drive

```
import torchvision
import glob
from matplotlib import pyplot as plt
import cv2
from google.colab.patches import cv2_imshow
import os.path as osp
import numpy as np
import torch
from torchvision import transforms, datasets, utils
from torch.utils.data import Dataset, DataLoader
from torchvision.models import vgg16
from torch.optim import lr_scheduler
import torch.optim as optim
import time
import copy
import torch.nn as nn
import PIL
import torch.nn.functional as Func
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

#Split the dataset into Train, Validation and Test set
total_image = sorted(glob.glob('drive/My Drive/CSCI-677-HW5/data_semantics/training/images/*'))
total_label = sorted(glob.glob('drive/My Drive/CSCI-677-HW5/data_semantics/training/labels/*'))
image_train, image_rest, label_train, label_rest = train_test_split(total_image, total_label, test_size=0.2)
image_test, image_val, label_test, label_val = train_test_split(image_rest, label_rest, test_size=0.5)

class semantic_dataset_train(Dataset):
    def __init__(self, split = 'test', transform = None):
        self void_labels = [0, 1, 2, 3, 4, 5, 6, 9, 10, 14, 15, 16, 18, 29, 30, -1]
        self valid_labels = [7, 8, 11, 12, 13, 17, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28]
        self ignore_index = 250
        self class_map = dict(zip(self.valid_labels, range(19)))
        self.split = split

        self.transform = transform

        self.img_list = image_train
        self.mask_list = label_train

    def __len__(self):
```

```

        return(len(self.img_list))

def __getitem__(self, idx):
    img = cv2.imread(self.img_list[idx])
    img = cv2.resize(img, (1242, 376))
    mask = None
    if self.split == 'train':
        mask = cv2.imread(self.mask_list[idx], cv2.IMREAD_GRAYSCALE)
        mask = cv2.resize(mask, (512, 256), interpolation=cv2.INTER_NEAREST)
        #mask = self.encode_segmap(mask)
        assert(mask.shape == (256, 512))

    if self.transform:
        img = self.transform(img)
        assert(img.shape == (3, 256, 512))
    else :
        assert(img.shape == (256, 512, 3))

    if self.split == 'train':
        return img, mask
    else :
        return img

class semantic_dataset_validation(Dataset):
    def __init__(self, split = 'test', transform = None):
        self.void_labels = [0, 1, 2, 3, 4, 5, 6, 9, 10, 14, 15, 16, 18, 29, 30, -1]
        self.valid_labels = [7, 8, 11, 12, 13, 17, 19, 20, 21, 22, 23, 24, 25, 26, 27,
        self.ignore_index = 250
        self.class_map = dict(zip(self.valid_labels, range(19)))
        self.split = split

        self.transform = transform

        self.img_list = image_val
        self.mask_list = label_val

    def __len__(self):
        return(len(self.img_list))

    def __getitem__(self, idx):
        img = cv2.imread(self.img_list[idx])
        img = cv2.resize(img, (1242, 376))
        mask = None
        if self.split == 'train':
            mask = cv2.imread(self.mask_list[idx], cv2.IMREAD_GRAYSCALE)
            mask = cv2.resize(mask, (512, 256), interpolation=cv2.INTER_NEAREST)
            #mask = self.encode_segmap(mask)
            assert(mask.shape == (256, 512))

        if self.transform:
            img = self.transform(img)

```

```

        assert(img.shape == (3, 256, 512))
    else :
        assert(img.shape == (256, 512, 3))

    if self.split == 'train':
        return img, mask
    else :
        return img

class semantic_dataset_test(Dataset):
    def __init__(self, split = 'test', transform = None):
        self.void_labels = [0, 1, 2, 3, 4, 5, 6, 9, 10, 14, 15, 16, 18, 29, 30, -1]
        self.valid_labels = [7, 8, 11, 12, 13, 17, 19, 20, 21, 22, 23, 24, 25, 26, 27,
        self.ignore_index = 250
        self.class_map = dict(zip(self.valid_labels, range(19)))
        self.split = split

        self.transform = transform

        self.img_list = image_test
        self.mask_list = label_test

    def __len__(self):
        return(len(self.img_list))

    def __getitem__(self, idx):
        img = cv2.imread(self.img_list[idx])
        img = cv2.resize(img, (1242, 376))
        mask = None
        if self.split == 'train':
            mask = cv2.imread(self.mask_list[idx], cv2.IMREAD_GRAYSCALE)
            mask = cv2.resize(mask, (512, 256), interpolation=cv2.INTER_NEAREST)
            #mask = self.encode_segmap(mask)
            assert(mask.shape == (256, 512))

        if self.transform:
            img = self.transform(img)
            assert(img.shape == (3, 256, 512))
        else :
            assert(img.shape == (256, 512, 3))

        if self.split == 'train':
            return img, mask
        else :
            return img

transformation = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize((256,512),interpolation=PIL.Image.BICUBIC),
    transforms.ToTensor()])

```

```

        transforms.Normalize(mean = [ 0.35675976, 0.37380189, 0.3764753], std = [0.
    ])
trainset = semantic_dataset_train(split = 'train', transform = transformation)
validationset = semantic_dataset_validation(split = 'train', transform = transformation)
testset = semantic_dataset_test(split = 'train', transform = transformation)

trainImages = DataLoader(trainset, batch_size = 5, shuffle = True)
valImages = DataLoader(validationset, batch_size = 5, shuffle = True)
testImages = DataLoader(testset, batch_size = 1, shuffle = True)

#FC32 architecture
vgg16_mod = vgg16(pretrained=True)
for param in vgg16_mod.features.parameters():
    param.requires_grad = False
class FCN32(nn.Module):
    def __init__(self):
        super(FCN32, self).__init__()
        self.features = vgg16_mod.features
        self.classifier = nn.Sequential(
            nn.Conv2d(512, 4096, 7),
            nn.ReLU(inplace=True),
            nn.Dropout2d(),
            nn.Conv2d(4096, 4096, 1),
            nn.ReLU(inplace=True),
            nn.Dropout2d(),
            nn.Conv2d(4096, 35, 1),
            nn.ConvTranspose2d(35, 35, 224, stride=32)
        )
    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x

fcn = FCN32()
print(fcn)
input = torch.randn(5,3,256,512)
output = fcn(input)
print(output.shape)

#FC16 architecture
vgg16_mod = vgg16(pretrained=True)
for param in vgg16_mod.features.parameters():
    param.requires_grad = False
class FCN16(nn.Module):
    def __init__(self):
        super(FCN16, self).__init__()
        self.features = vgg16_mod.features
        self.classifier = nn.Sequential(

```

```

        nn.Conv2d(512, 4096, 7),
        nn.ReLU(inplace=True),
        nn.Conv2d(4096, 4096, 1),
        nn.ReLU(inplace=True),
        nn.Conv2d(4096, 35, 1)
    )
    self.score_pool4 = nn.Conv2d(512, 35, 1)
    self.upscore2 = nn.ConvTranspose2d(35, 35, 14, stride=2, bias=False)
    self.upscore16 = nn.ConvTranspose2d(35, 35, 16, stride=16, bias=False)

def forward(self, x):
    pool4 = self.features[:-7](x)
    pool5 = self.features[-7:](pool4)
    pool5_upscored = self.upscore2(self.classifier(pool5))
    pool4_scored = self.score_pool4(pool4)
    combined = pool4_scored + pool5_upscored
    res = self.upscore16(combined)
    return res

fcn16 = FCN16()
print(fcn16)
input = torch.randn(5, 3, 256, 512)
output = fcn16(input)
print(output.shape)

FCN16(
    features: Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU(inplace=True)
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (6): ReLU(inplace=True)
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (8): ReLU(inplace=True)
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace=True)
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (13): ReLU(inplace=True)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (18): ReLU(inplace=True)
        (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (20): ReLU(inplace=True)
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (22): ReLU(inplace=True)
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (25): ReLU(inplace=True)
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
)
```

```

(27): ReLU(inplace=True)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Conv2d(512, 4096, kernel_size=(7, 7), stride=(1, 1))
(1): ReLU(inplace=True)
(2): Conv2d(4096, 4096, kernel_size=(1, 1), stride=(1, 1))
(3): ReLU(inplace=True)
(4): Conv2d(4096, 35, kernel_size=(1, 1), stride=(1, 1))
)
(score_pool4): Conv2d(512, 35, kernel_size=(1, 1), stride=(1, 1))
(upscore2): ConvTranspose2d(35, 35, kernel_size=(14, 14), stride=(2, 2), bias=False)
(upscore16): ConvTranspose2d(35, 35, kernel_size=(16, 16), stride=(16, 16), bias=False)
)
torch.Size([5, 35, 256, 512])

```

Double-click (or enter) to edit

```

#Do not use
def training(fcn, criterion, optimizer,num_epochs):

    val_loss = []
    train_loss = []
    for epoch in range(num_epochs):
        fcn.train()
        val_running_loss = 0
        train_running_loss = 0
        fcn16.train()
        for i, trainData in enumerate(trainImages, 0):
            imgs, labels = trainData
            labels = labels.type(torch.LongTensor)
            #set the gradient to zero
            optimizer.zero_grad()
            outputs = fcn(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            train_running_loss += loss.item()
            optimizer.step()

        train_loss.append(train_running_loss)
        fcn16.eval()
        with torch.no_grad():
            for k, dat1 in enumerate(valImages):
                # Get image, label pair
                inputs1, labels1 = dat1
                labels1 = labels1.type(torch.LongTensor)
                # Predicting segmentation for val inputs
                outputs1 = fcn(inputs1)

                # Compute CE loss and aggregate it

```

```

        loss1 = criterion(outputs1, labels1)
        val_running_loss += loss1.item()

    val_loss.append(val_running_loss)
    print("Epochs: ", epoch, "Training Loss: ", train_running_loss, "val_running_loss: "
return train_loss, val_loss

criterion = nn.CrossEntropyLoss()

```

```

optimizer = optim.Adam(fcn16.parameters(), lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_decay=0)
train_loss, val_loss = training(fcn16, criterion, optimizer, num_epochs=40)
torch.save(fcn16, 'drive/My Drive/CSCI-677-HW5/data_semantics/ModelFCN16_Adam_dropout.pt')

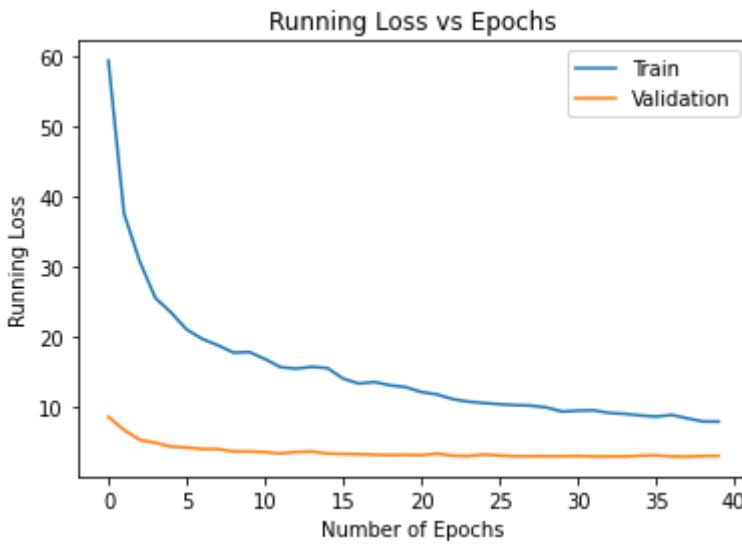
```

Epochs:	Training Loss:	val_running_loss:
0	59.28895115852356	8.55336415767669
1	37.511099100112915	6.6915805339813
2	30.72277069091797	5.29584228992462
3	25.50015300512314	4.85788941383361
4	23.424445688724518	4.32322156429291
5	21.01056444644928	4.19649475812912
6	19.65706741809845	3.97957623004913
7	18.750891357660294	3.9666209816932
8	17.704172015190125	3.6421911120414
9	17.79672572016716	3.62778818607330
10	16.8015516102314	3.53162536025047
11	15.630626678466797	3.347768396139
12	15.435338705778122	3.562463164329
13	15.679339826107025	3.624943733215
14	15.511705696582794	3.338151812553
15	14.003555357456207	3.282106667757
16	13.2979174554348	3.24957343935966
17	13.527827441692352	3.164175301790
18	13.056245625019073	3.111047983169
19	12.79800832271576	3.1354165673255
20	12.091058164834976	3.098386257886
21	11.756760269403458	3.323573887348
22	11.086442202329636	3.029297202825
23	10.726012796163559	2.997289389371
24	10.530326664447784	3.186679244041
25	10.360094755887985	3.054947137832
26	10.235976576805115	2.953990906476
27	10.16593188047409	2.9485259652137
28	9.887732744216919	2.9563983082771
29	9.317112028598785	2.9550115168094
30	9.44376426935196	2.99237030744552
31	9.49117062985897	2.90960735082626
32	9.119388341903687	2.9272747337818
33	8.997102707624435	2.9211678802967
34	8.75710615515709	3.02060854434967
35	8.606159165501595	3.0815064311027
36	8.827051058411598	2.9174649417400
37	8.356737405061722	2.8964205086231
38	7.899606853723526	2.9710412323474
39	7.882933139801025	3.0042795538902

```
#torch.save(fcn, 'drive/My Drive/CSCI677-Homeworks/data_semantics/Model_3_val_SGD.pt')
```

```
#Plotting Running loss vs Epochs
epochList = [i for i in range(0, 40)]
plt.plot(epochList, [x for x in train_loss], label = "Train")
plt.plot(epochList, [y for y in val_loss], label = "Validation")

plt.xlabel('Number of Epochs')
plt.ylabel('Running Loss')
plt.title('Running Loss vs Epochs')
plt.legend()
plt.show()
```



```
#Calculating IoU and mean IoU
def get_mean_iou(conf_mat, multiplier=1.0):
    cm = conf_mat.copy()
    np.fill_diagonal(cm, np.diag(cm) * multiplier)
    inter = np.diag(cm)
    gt_set = cm.sum(axis=1)
    pred_set = cm.sum(axis=0)
    union_set = gt_set + pred_set - inter
    iou = inter.astype(float) / union_set
    mean_iou = np.nanmean(iou)
    return mean_iou

overall_conf_mat = np.zeros((35, 35))
with torch.no_grad():
    for k, dat1 in enumerate(testImages):
        # Get image, label pair
        inputs1, labels1 = dat1

        # Predicting segmentation for val inputs
        outputs1 = fcn16(inputs1)
```

```
preds = torch.argmax(outputs1, dim=1).detach().numpy()
gt = labels1.detach().numpy()

# Compute confusion matrix
conf_mat = confusion_matrix(y_pred=preds.flatten(), y_true=gt.flatten(), labels=labels)
overall_conf_mat += conf_mat

print('DICE score: {}'.format(np.round(get_mean_iou(conf_mat=overall_conf_mat, multiplier=1.0), 2)))
print('Mean IOU score: {}'.format(np.round(get_mean_iou(conf_mat=overall_conf_mat), 2)))

DICE score: 0.37
Mean IOU score: 0.28
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:9: RuntimeWarning: ...
  if __name__ == '__main__':


class UnNormalize(object):
    def __init__(self, mean, std):
        self.mean = mean
        self.std = std

    def __call__(self, tensor):
        """
        Args:
            tensor (Tensor): Tensor image of size (C, H, W) to be normalized.
        Returns:
            Tensor: Normalized image.
        """
        for t, m, s in zip(tensor, self.mean, self.std):
            t.mul_(s).add_(m)
            # The normalize code -> t.sub_(m).div_(s)
        return tensor

import numpy as np
#Map the labels to colors and display the results
def get_mean_iou(conf_mat, multiplier=1.0):
    cm = conf_mat.copy()
    np.fill_diagonal(cm, np.diag(cm) * multiplier)
    inter = np.diag(cm)
    gt_set = cm.sum(axis=1)
    pred_set = cm.sum(axis=0)
    union_set = gt_set + pred_set - inter
    iou = inter.astype(float) / union_set
    mean_iou = np.nanmean(iou)
    return mean_iou

def decode_segmap(image, nc=35):

    label_colors = np.array([( 0,  0,  0),
                            ( 0,  0,  0), ( 0,  0,  0), ( 0,  0,  0), ( 0,  0,  0), (111, 74,
                            122), (150, 100), ( 70,  70,  70), (102, 102, 156), (100, 152, 152), (100, 152, 152)]
```

```
FCN_semantic_segmentation.ipynb - Colaboratory
(250, 150, 140), (10, 10, 10), (102, 102, 102), (190, 250, 250), (180, 180, 180),
(250, 170, 30), (220, 220, 0), (107, 142, 35), (152, 251, 152), (70, 130, 180),
(0, 0, 90), (0, 0, 110), (0, 80, 100), (0, 0, 230), (119, 11, 11)
```

```
r = np.zeros_like(image).astype(np.uint8)
g = np.zeros_like(image).astype(np.uint8)
b = np.zeros_like(image).astype(np.uint8)

for l in range(0, nc):
    idx = image == l
    r[idx] = label_colors[l, 0]
    g[idx] = label_colors[l, 1]
    b[idx] = label_colors[l, 2]

rgb = np.stack([r, g, b], axis=2)
return rgb

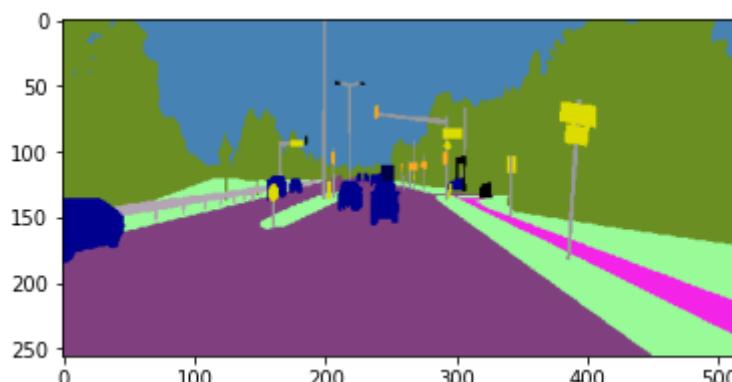
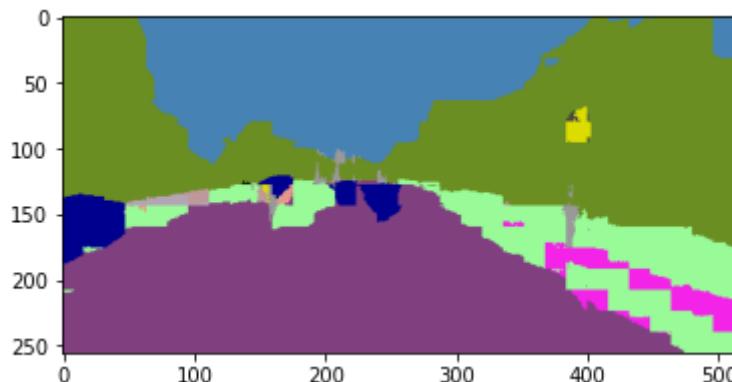
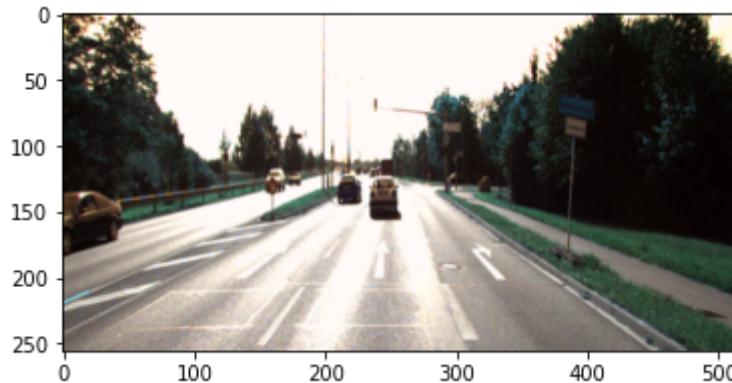
with torch.no_grad():
    for k, dat in enumerate(testImages):
        inputs, labels = dat
        print("Test set Number: ", k)
        outputs = fcn16(inputs)
        pred = torch.argmax(outputs[0].squeeze(), dim=0).detach().numpy()
        gt = labels[0]

        #Calculate IoU
        preds = torch.argmax(outputs, dim=1).detach().numpy()
        groundTruth = labels.detach().numpy()
        conf_mat = confusion_matrix(y_pred=preds.flatten(), y_true=groundTruth.flatten())
        IoU = get_mean_iou(conf_mat)
        Dice = get_mean_iou(conf_mat, multiplier=2.0)
        print("IoU: ", IoU)
        print("Dice: ", Dice)

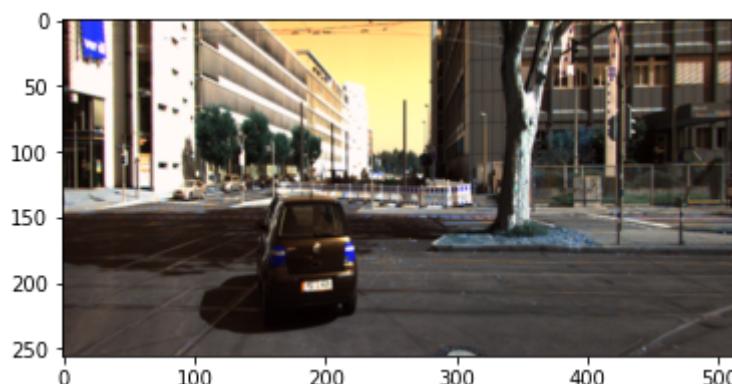
    unorm = UnNormalize(mean = [0.35675976, 0.37380189, 0.3764753], std = [0.32064945,
    img = unorm(inputs)

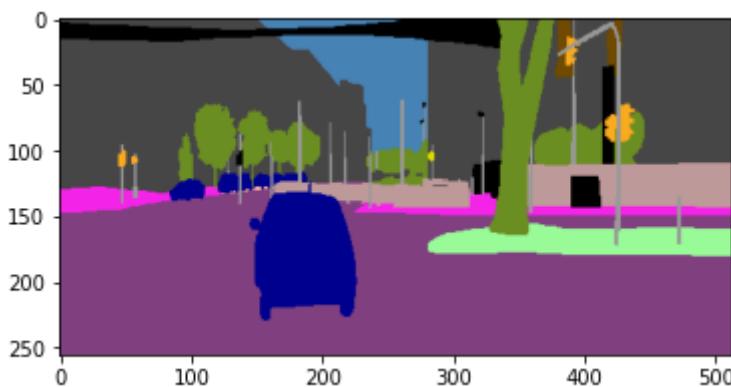
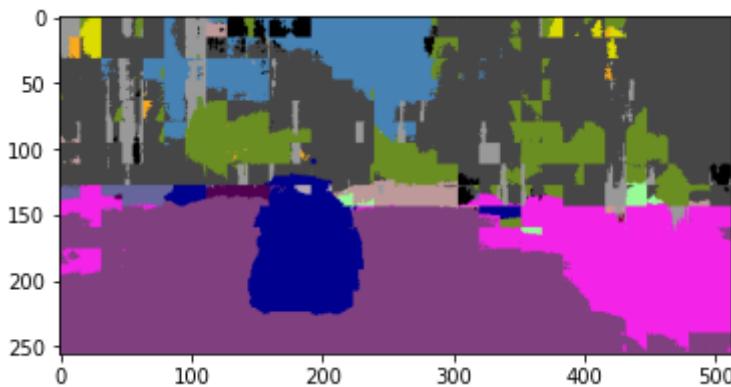
    plt.imshow(img[0].permute(1, 2, 0))
    plt.show()
    segmentation = decode_segmap(pred)
    gt_colored = decode_segmap(gt)
    plt.imshow(segmentation)
    plt.show()
    plt.imshow(gt_colored)
    plt.show()
```

```
Test set Number: 0
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:10: RuntimeWarning:
  # Remove the CWD from sys.path while we load stuff.
Clipping input data to the valid range for imshow with RGB data ([0..1] for float
IoU:  0.2994560619263632
Dice:  0.35667136170434527
```



```
Test set Number: 1
Clipping input data to the valid range for imshow with RGB data ([0..1] for float
IoU:  0.16395437102933613
Dice:  0.21569265167328563
```



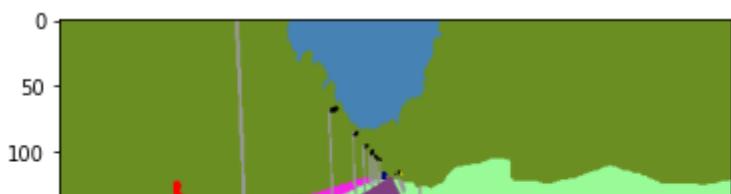
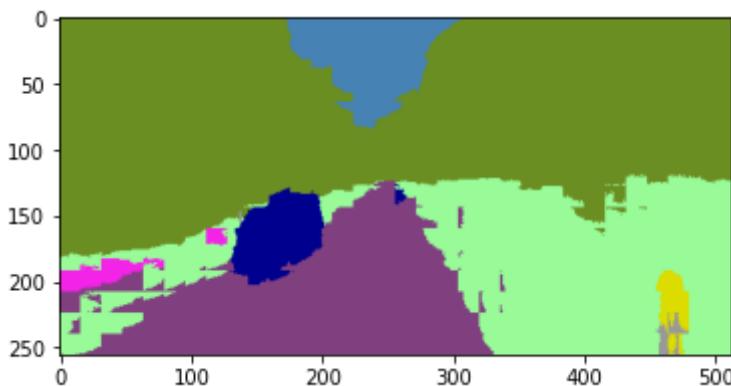
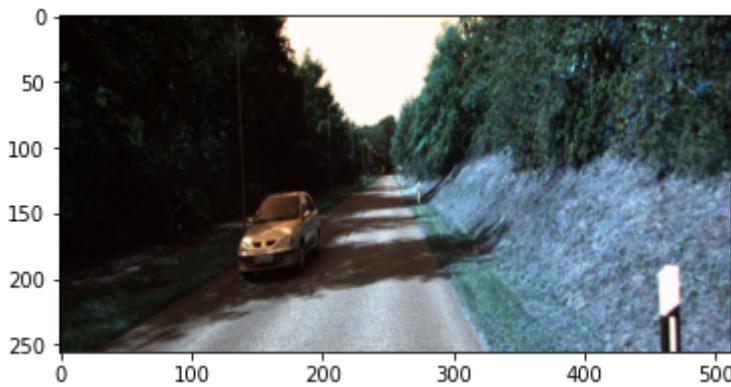


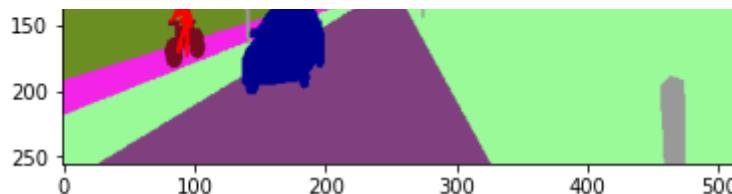
Test set Number: 2

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.4302027937450505

Dice: 0.4875077439789489



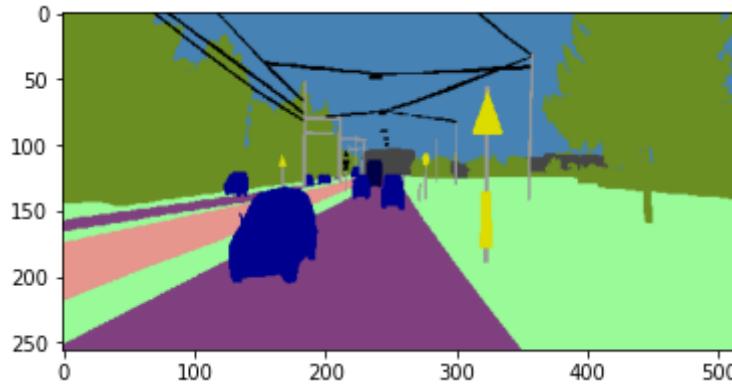
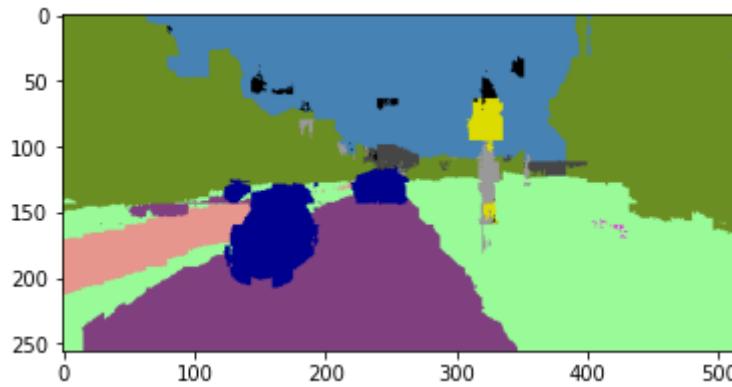
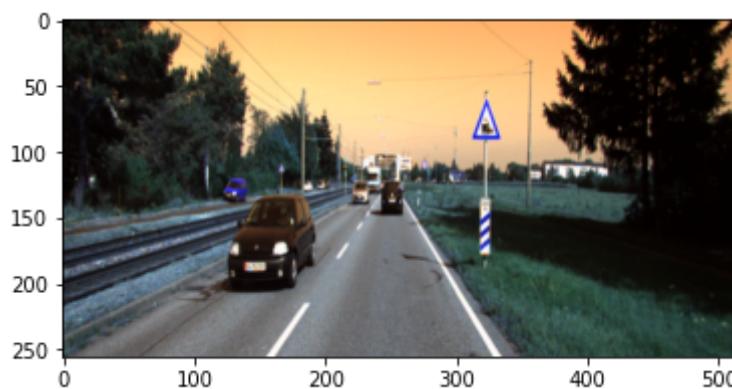


Test set Number: 3

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.480692582725838

Dice: 0.5609770772176007



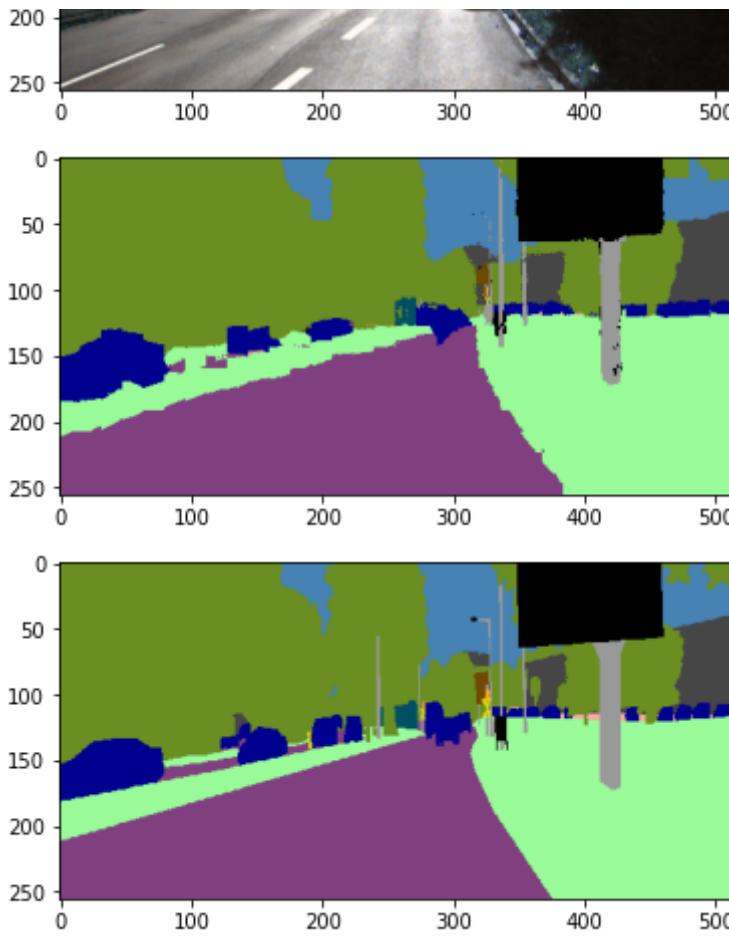
Test set Number: 4

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.558094593113312

Dice: 0.633653204353108



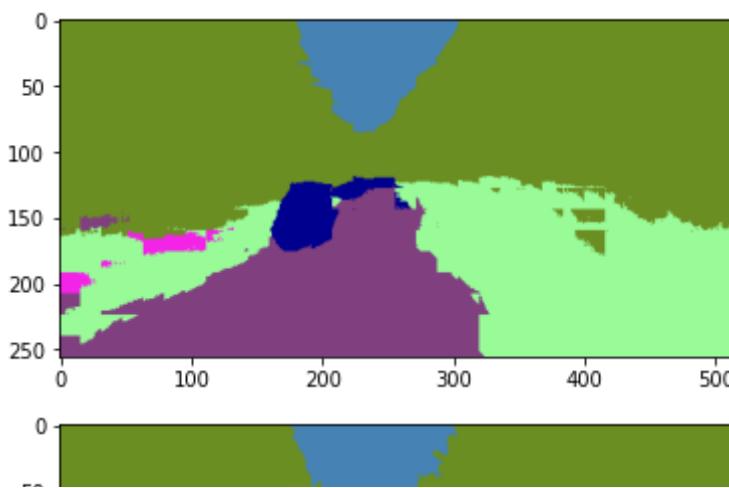
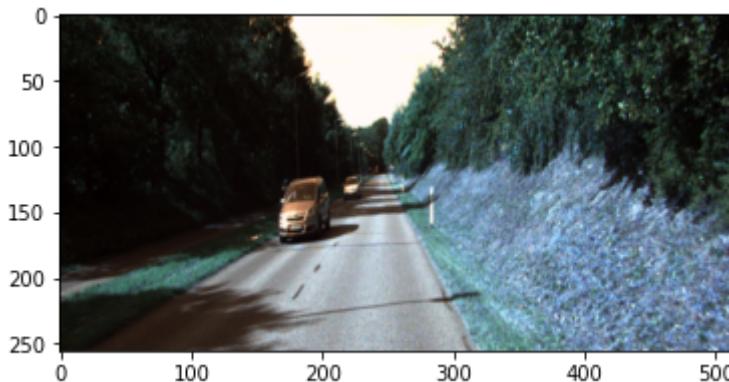


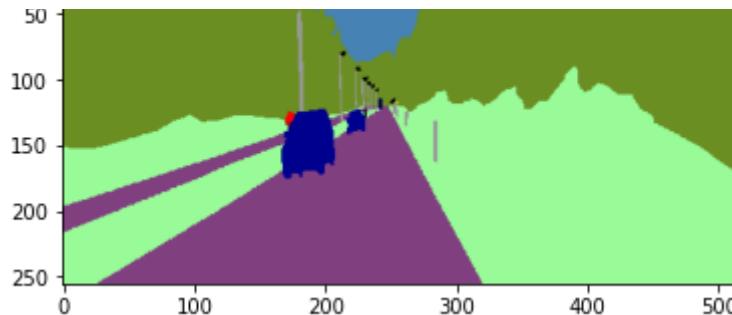
Test set Number: 5

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.32771281793950396

Dice: 0.3655345298864862



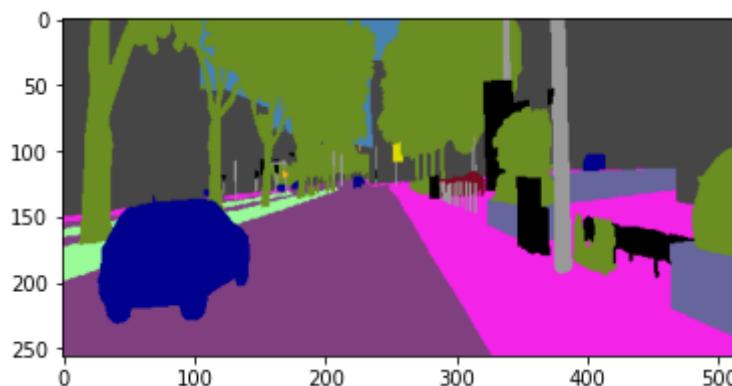
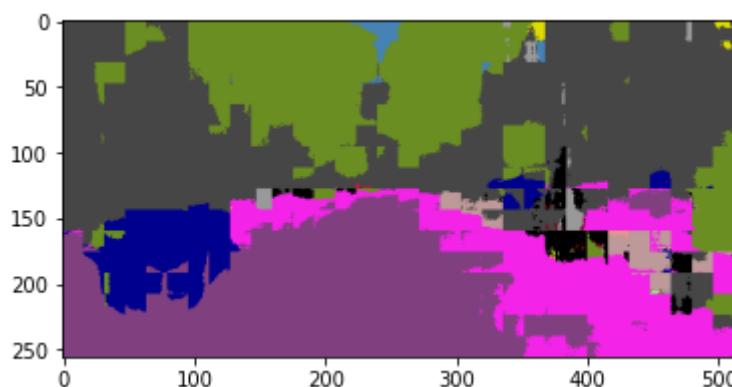
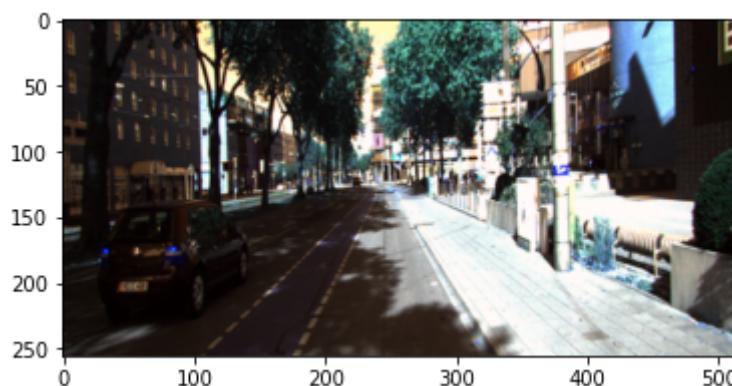


Test set Number: 6

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.20680618190206435

Dice: 0.2630828931842847

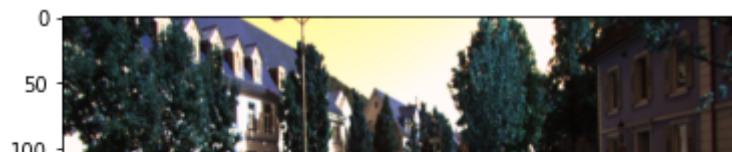


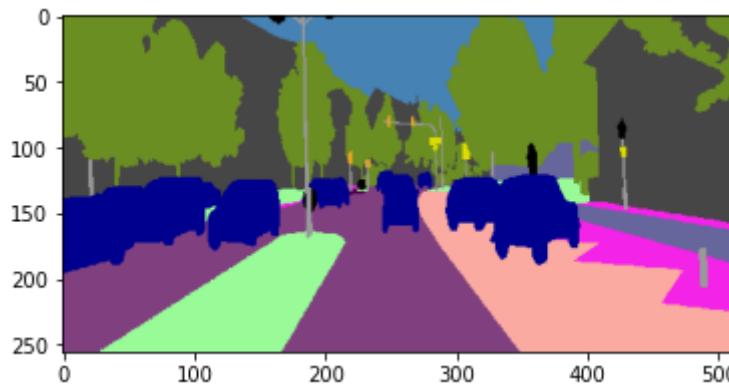
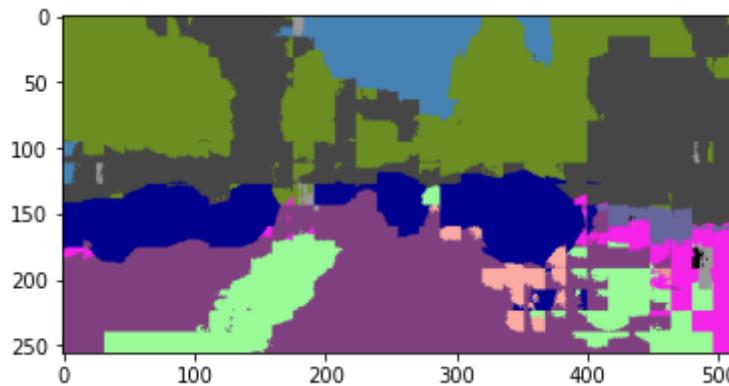
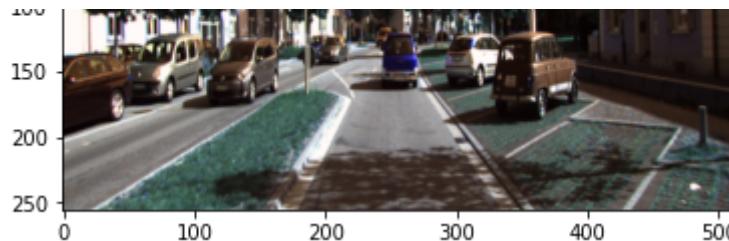
Test set Number: 7

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.30699851239735787

Dice: 0.38566679687995015



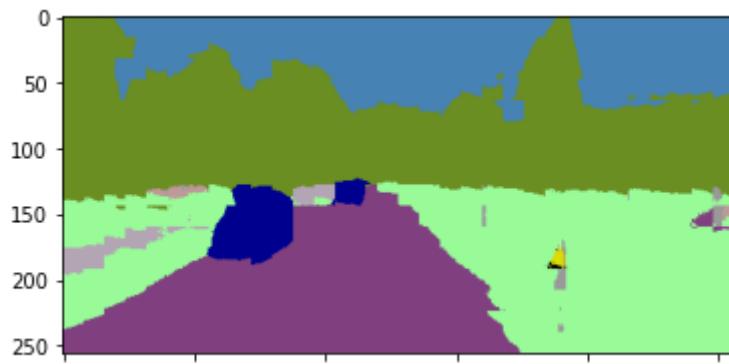
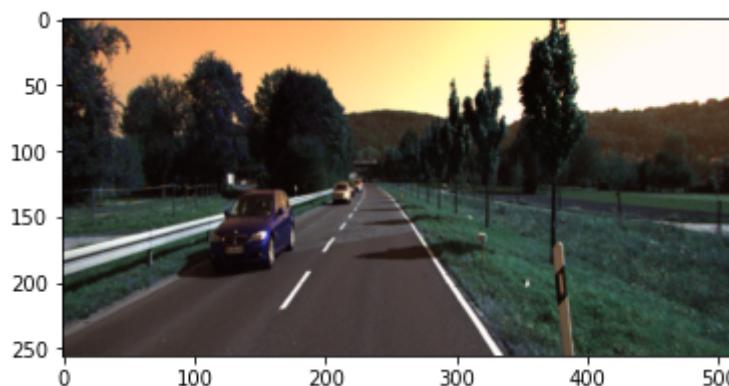


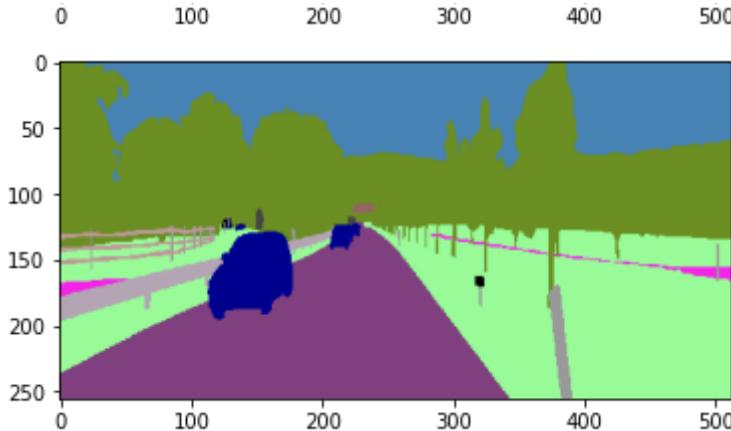
Test set Number: 8

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.36153677072828555

Dice: 0.4056605843911772



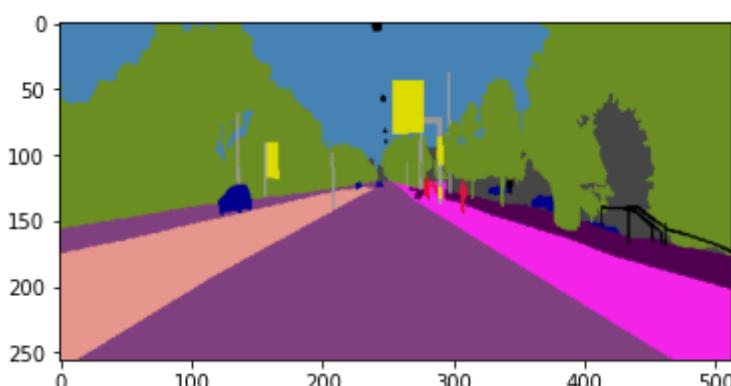
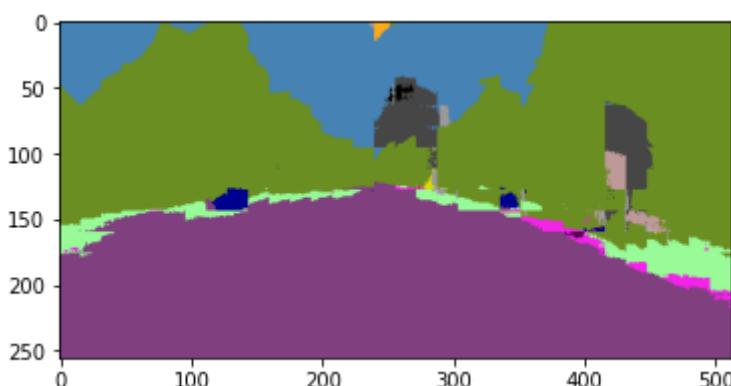
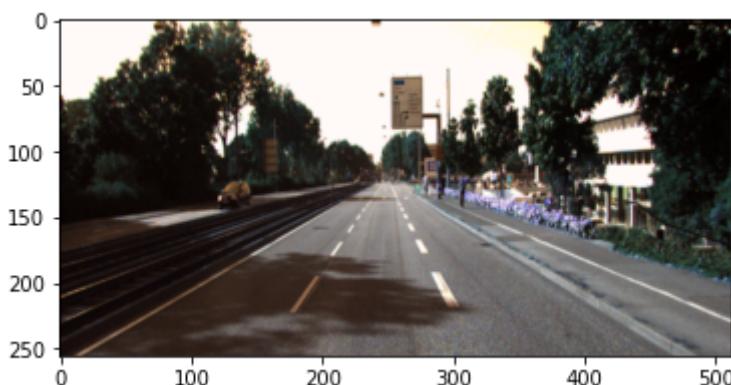


Test set Number: 9

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.16616832804771853

Dice: 0.20437211327839633



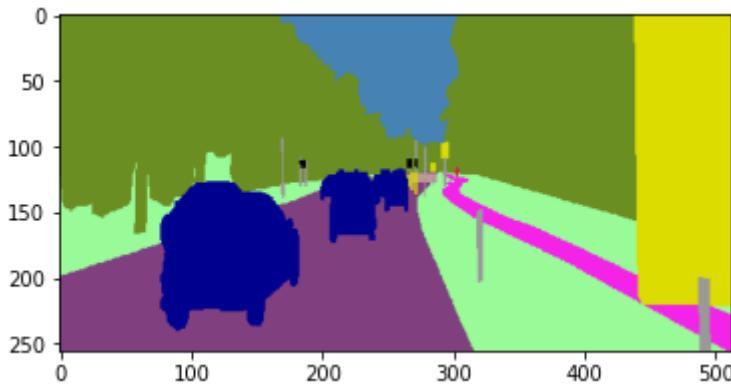
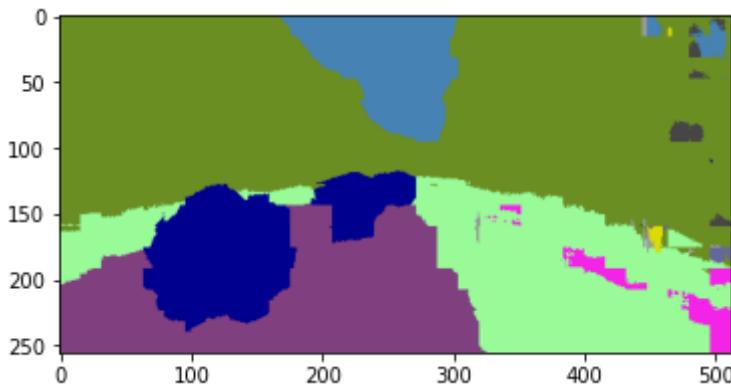
Test set Number: 10

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.30941439842409657

Dice: 0.35566833239236034



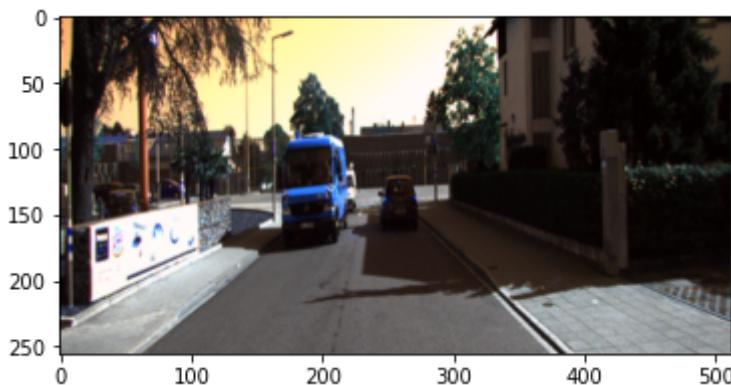


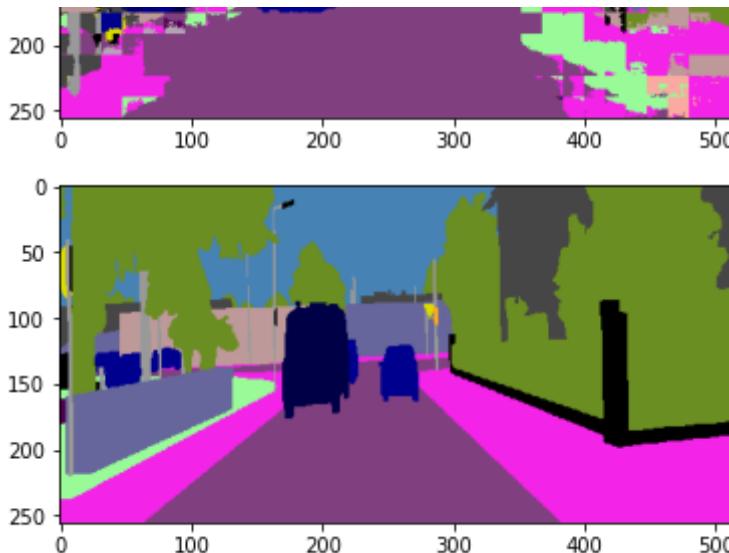
Test set Number: 11

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.16434791265206036

Dice: 0.213228998988609



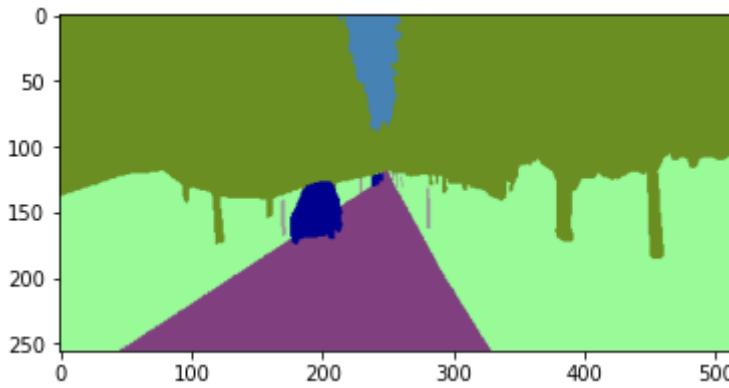
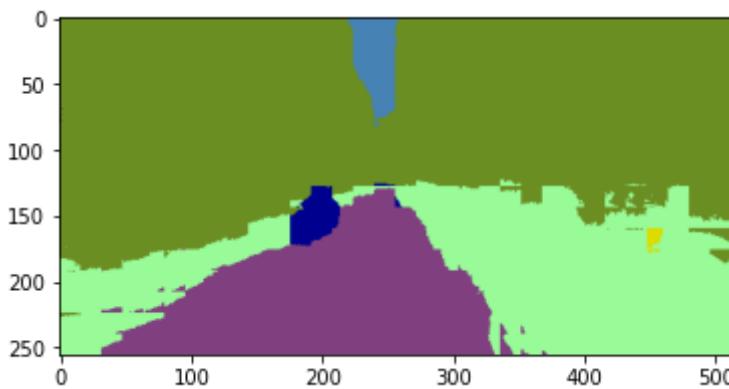
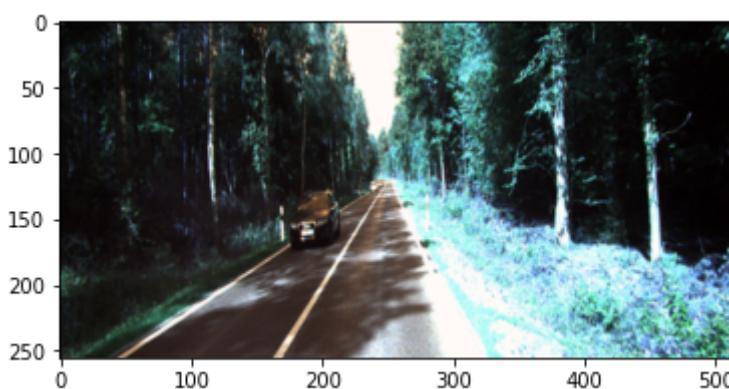


Test set Number: 12

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.387273178641238

Dice: 0.43556686956902546



Test set Number: 13

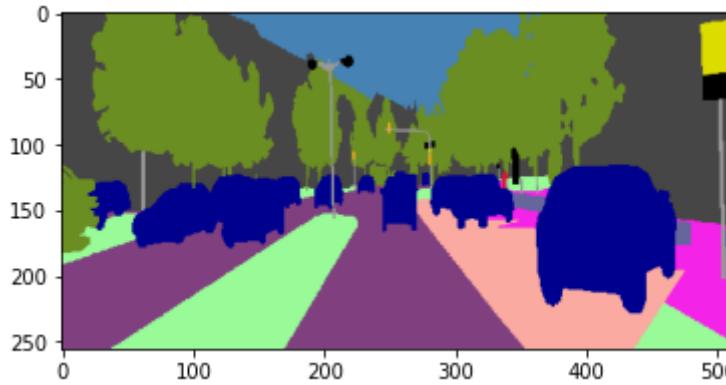
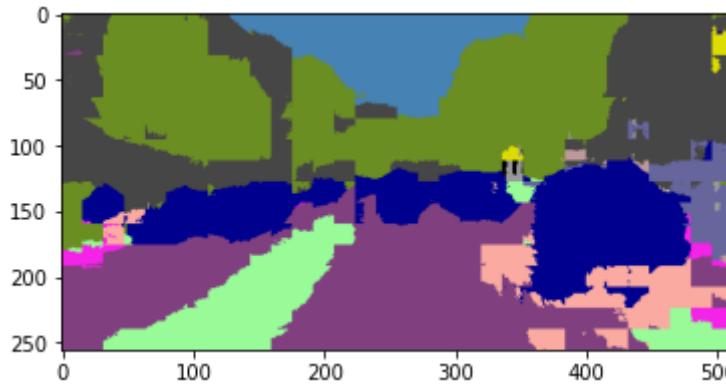
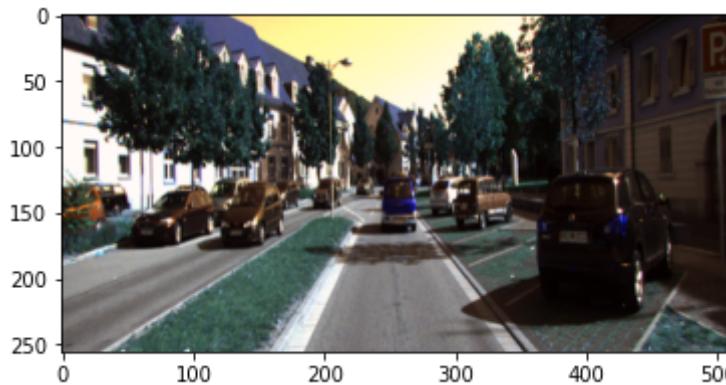
Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

<https://colab.research.google.com/drive/1zLwNCjzUIf3uxUmCHMz75oYcuredvI4Q?authuser=2#scrollTo=GU-2CEHhECUk&printMode=true>

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.33174588345034284

Dice: 0.4039578869473812

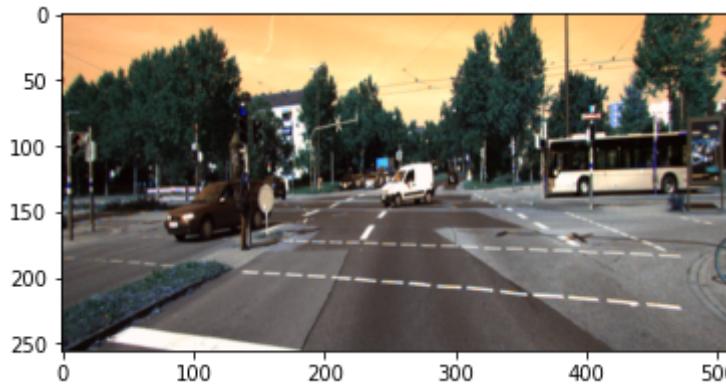


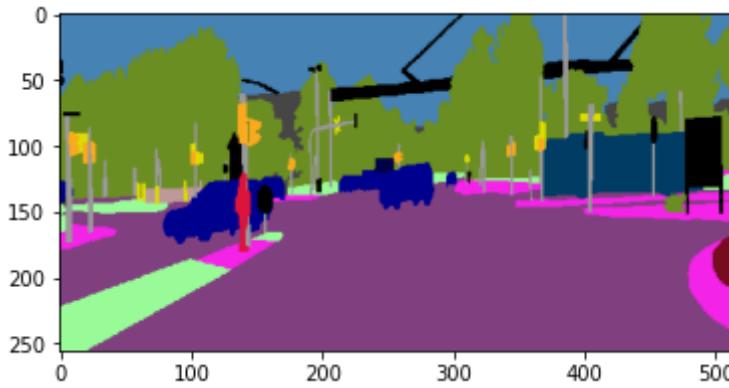
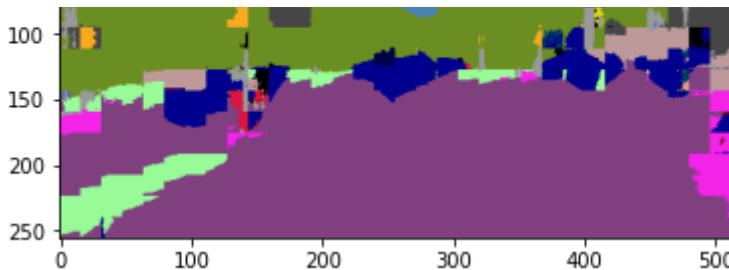
Test set Number: 14

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.24364473667441405

Dice: 0.32128596644451046



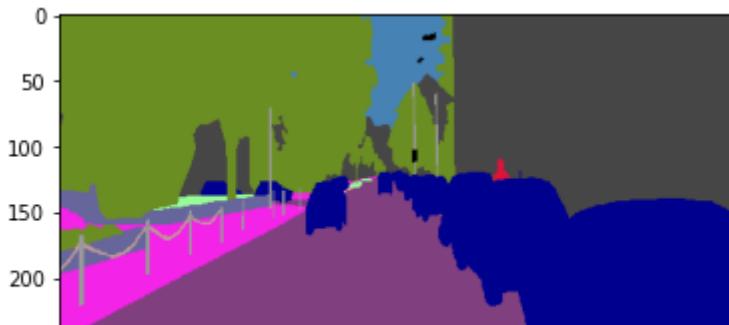
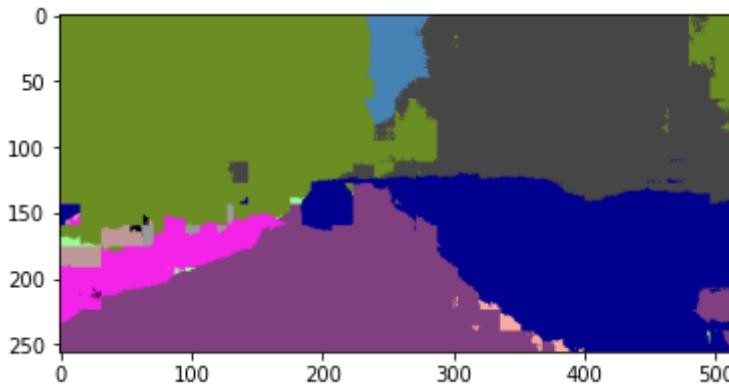
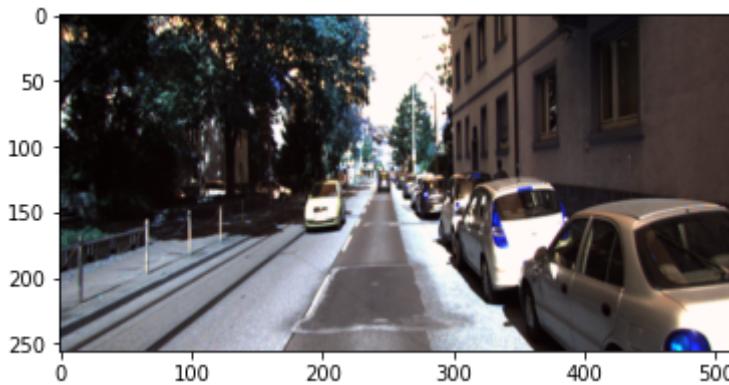


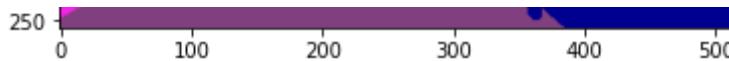
Test set Number: 15

Clipping input data to the valid range for imshow with RGB data ([0..1] for float)

IoU: 0.3371714156819543

Dice: 0.3889251375256572



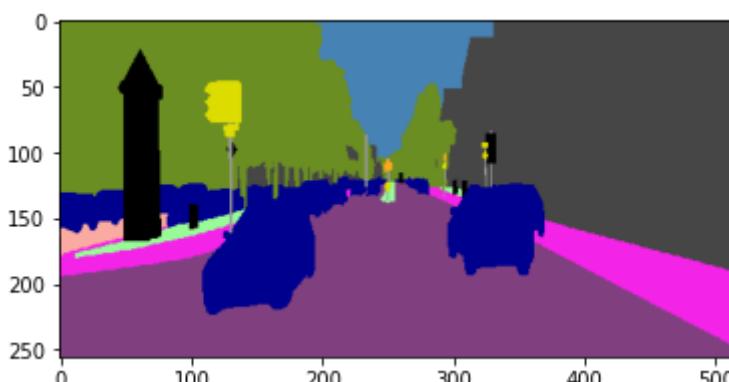
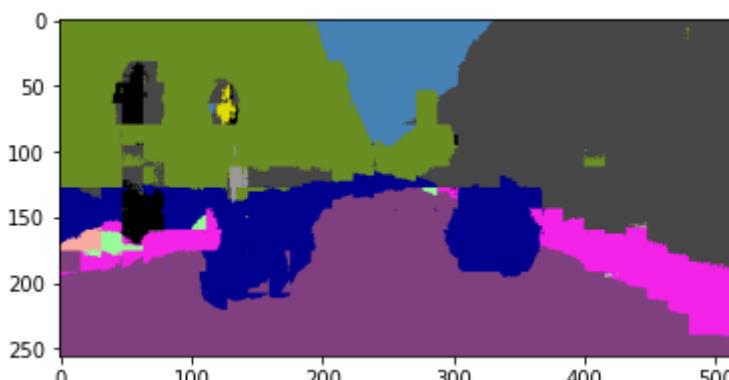


Test set Number: 16

Clipping input data to the valid range for imshow with RGB data ([0..1] for float32 data)

IoU: 0.44335624004307006

Dice: 0.5214470961050112

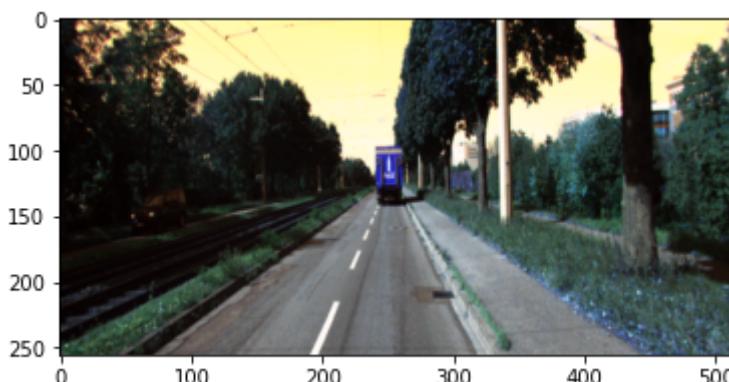


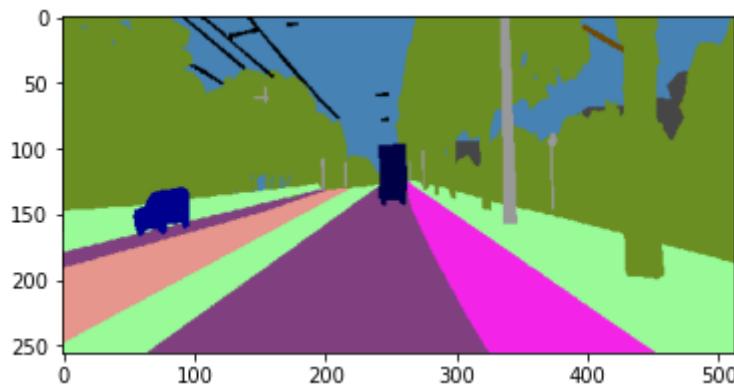
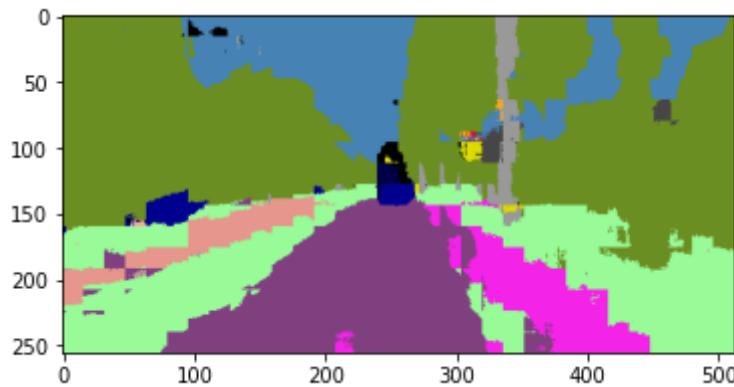
Test set Number: 17

Clipping input data to the valid range for imshow with RGB data ([0..1] for float32 data)

IoU: 0.26893518653329057

Dice: 0.340992803263865

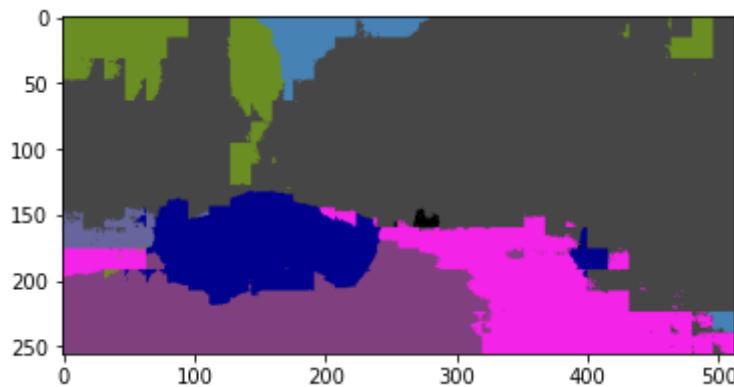


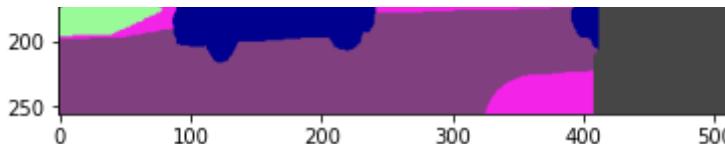


Test set Number: 18

IoU: 0.33603373597090974

Dice: 0.4106748131969363



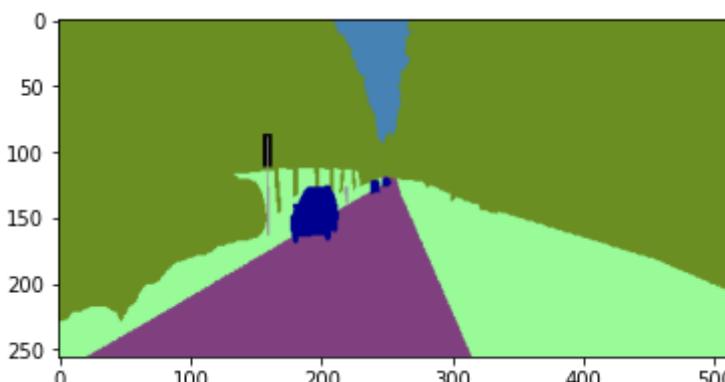
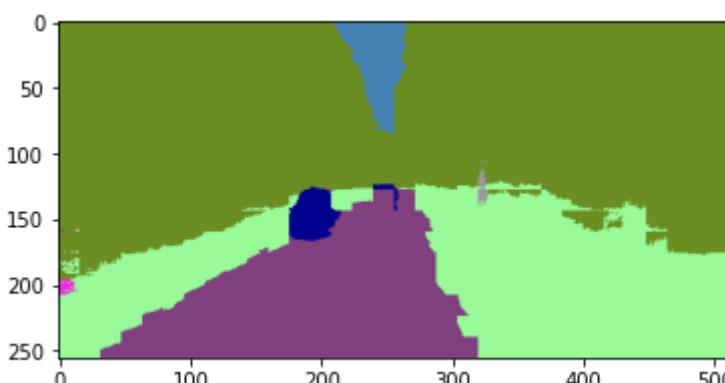
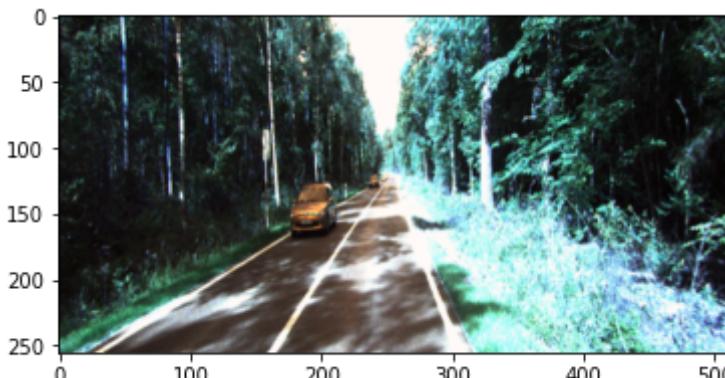


Test set Number: 19

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.45357700333428064

Dice: 0.497912989243278



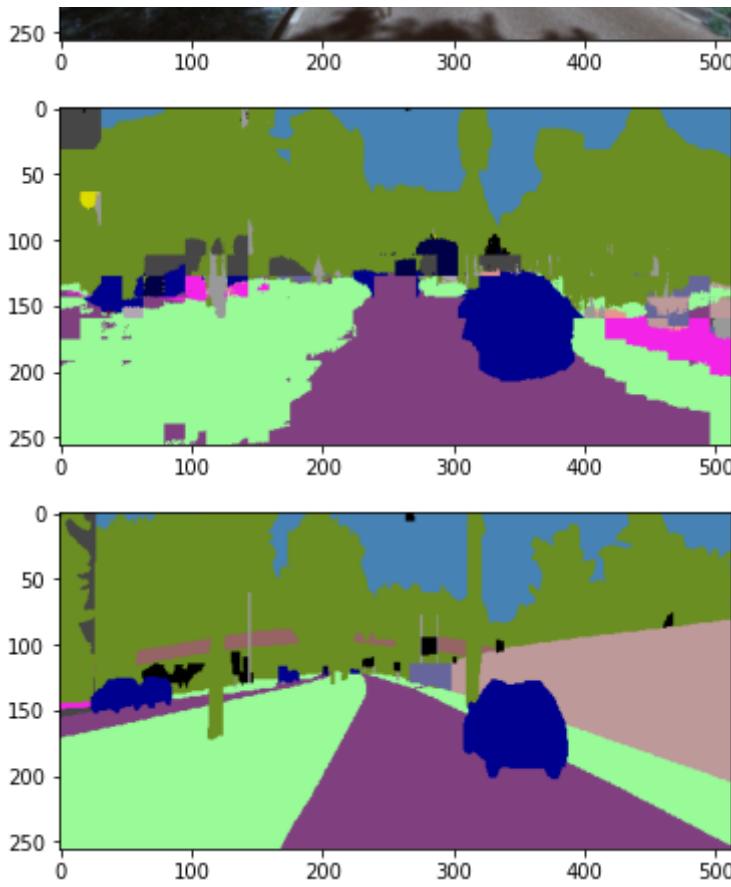
Test set Number: 20

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.22430682276647657

Dice: 0.2627560128761261



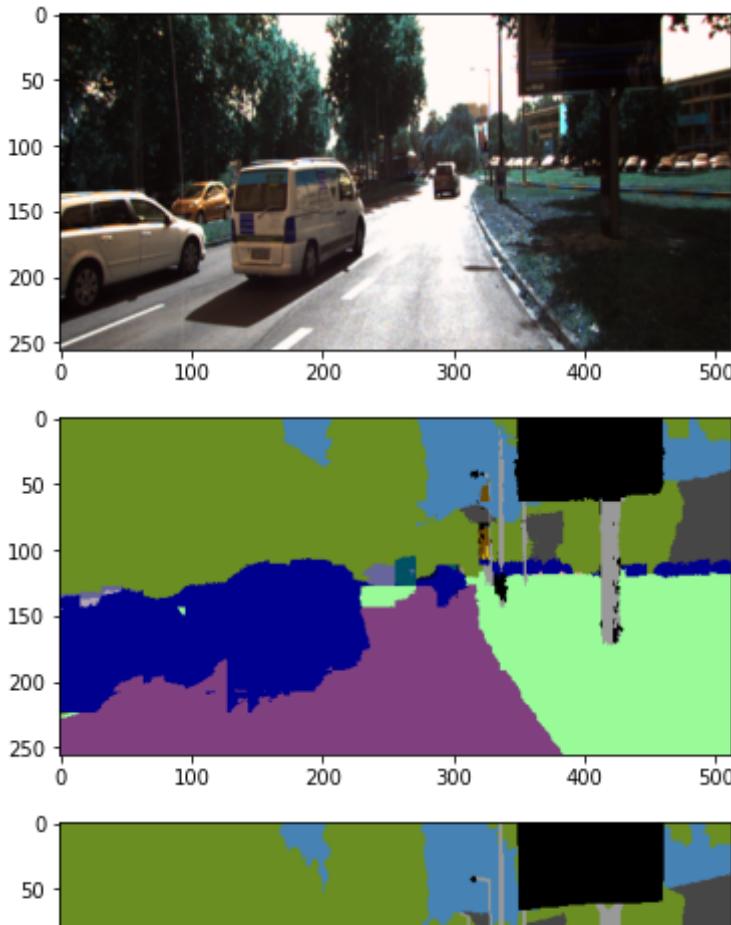


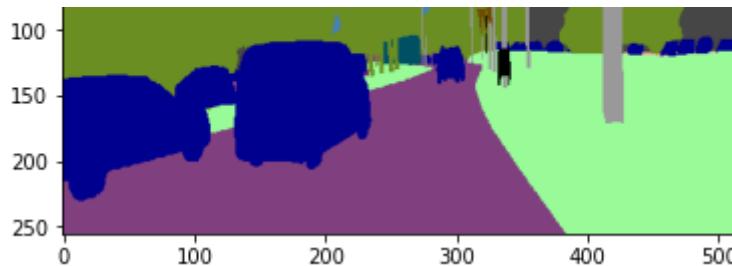
Test set Number: 21

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.49505334557847913

Dice: 0.5452795421055631



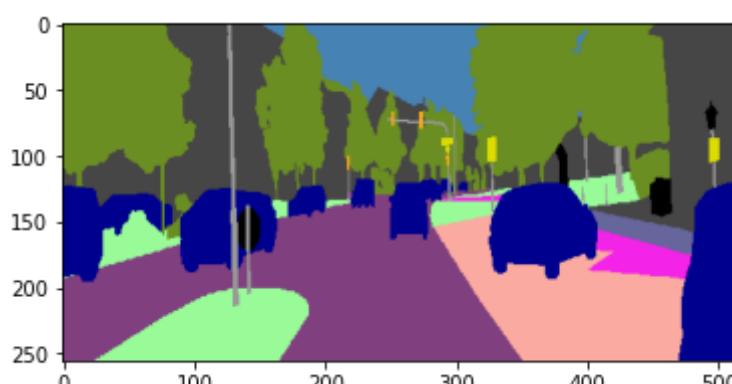
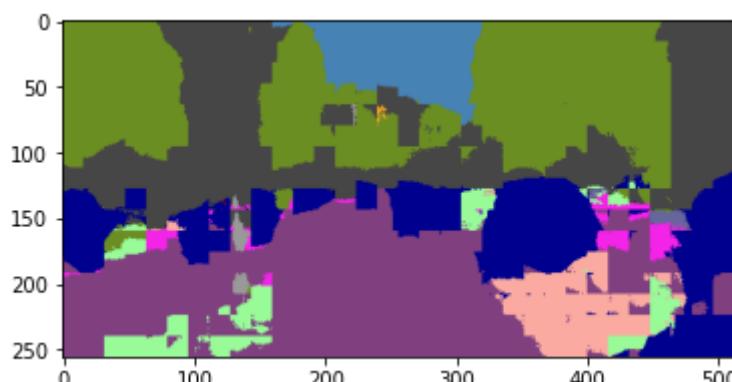


Test set Number: 22

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.27982476478884394

Dice: 0.3451869345183699

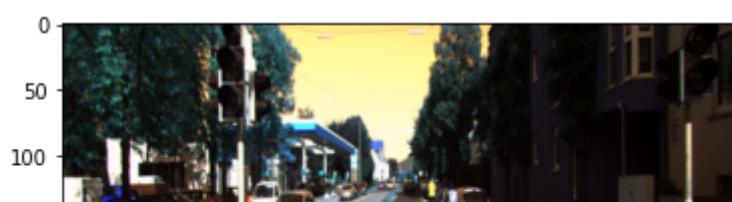


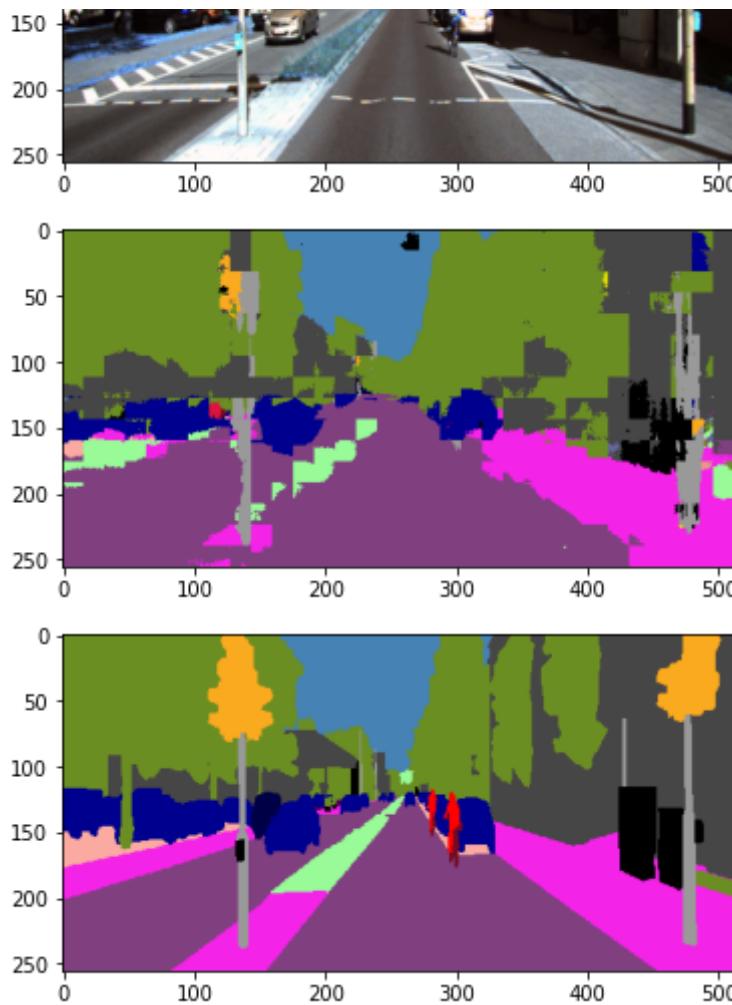
Test set Number: 23

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.2538665278997274

Dice: 0.32608584871224455



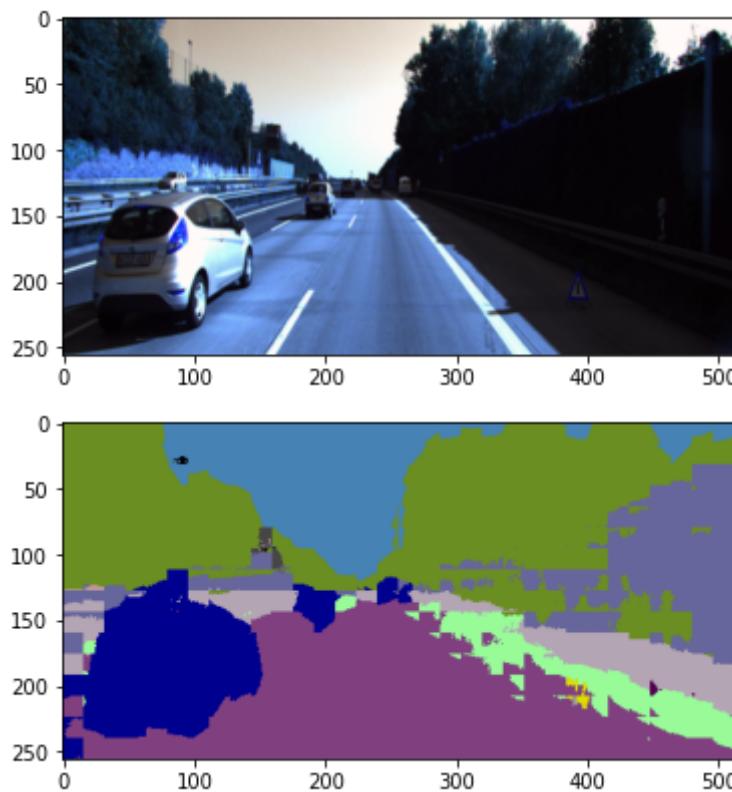


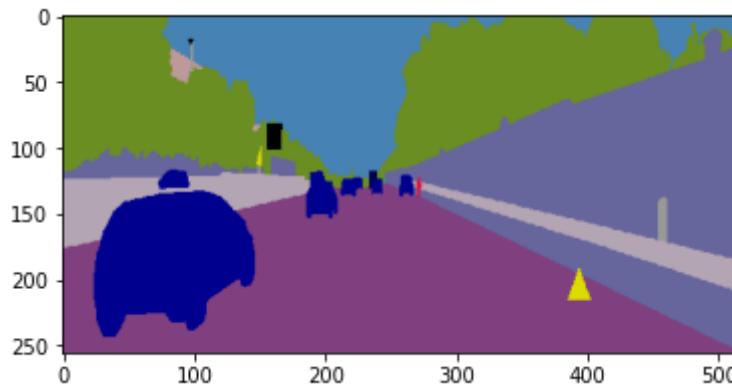
Test set Number: 24

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.2850189576481091

Dice: 0.342543047549985



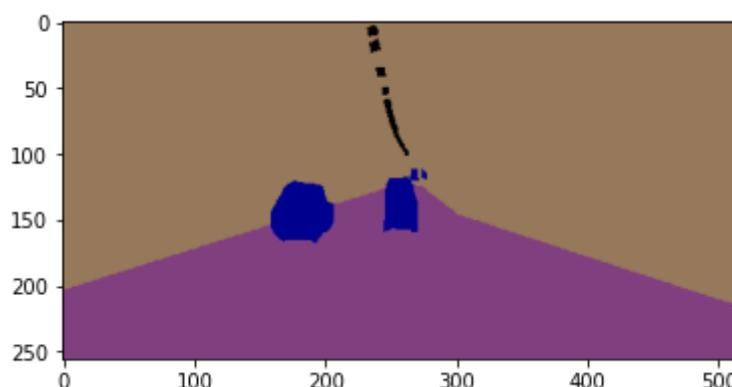
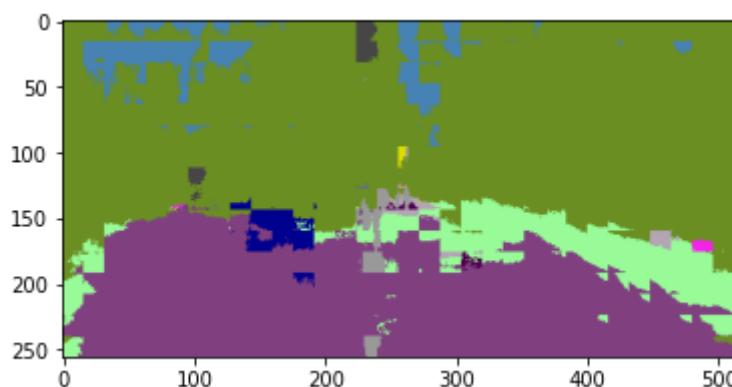
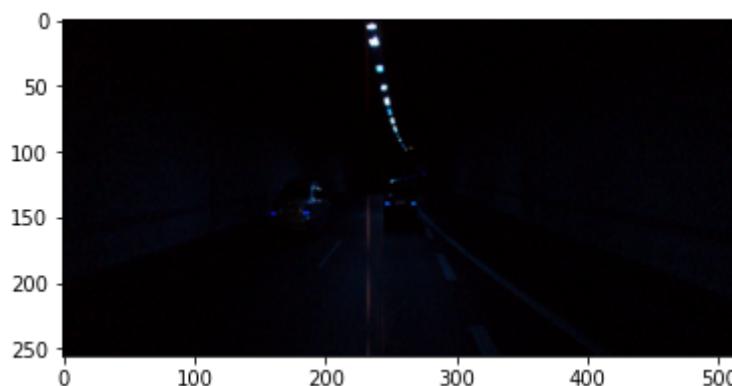


Test set Number: 25

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.06177069300187014

Dice: 0.07592515196681485

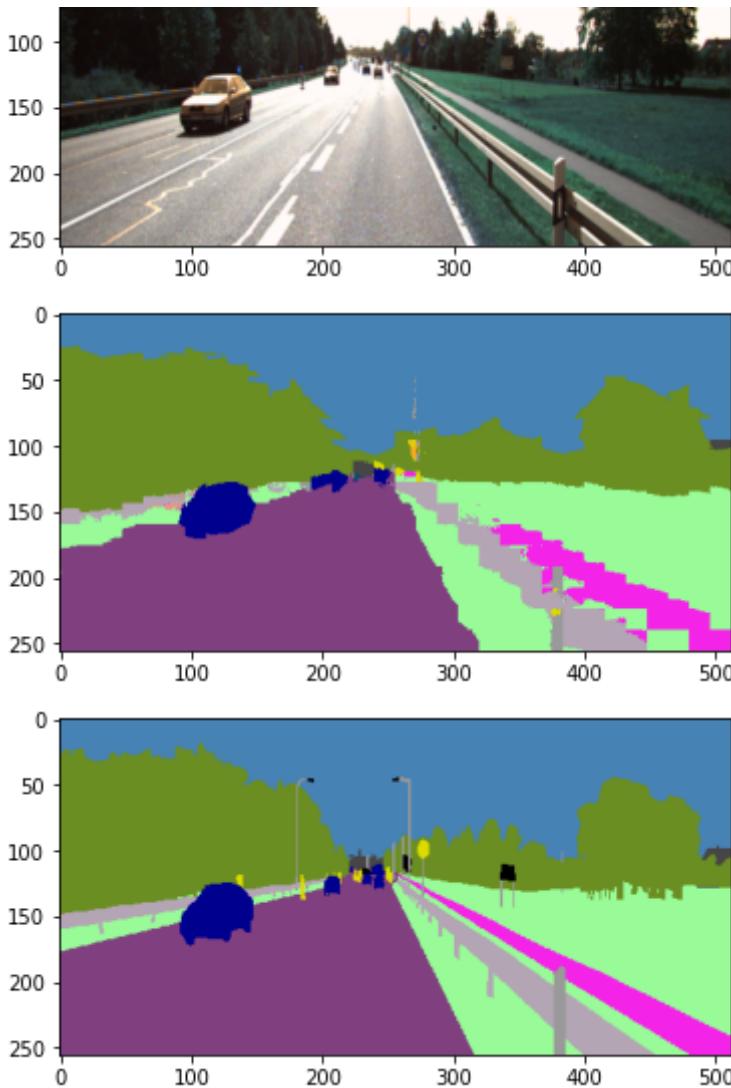


Test set Number: 26

IoU: 0.3742639293547273

Dice: 0.4288261269461293



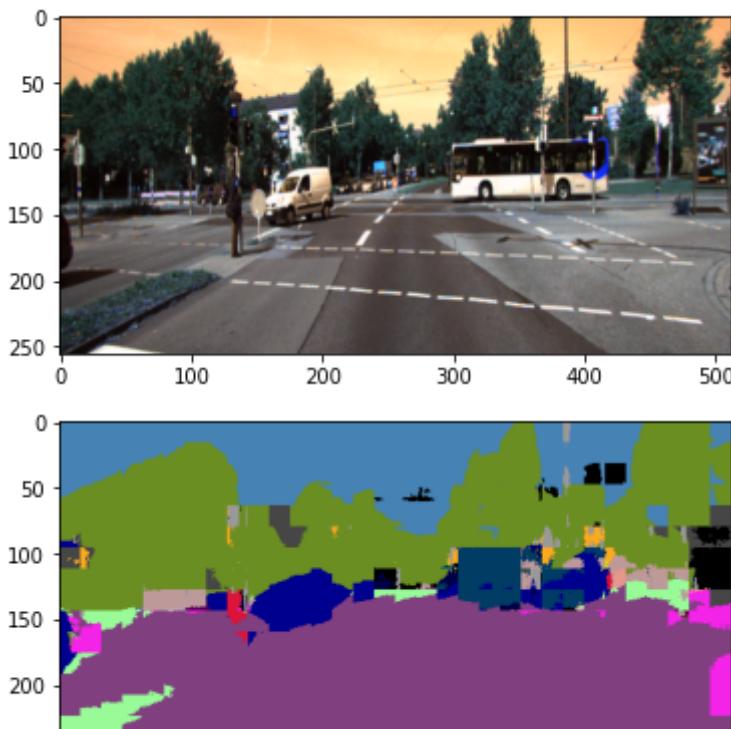


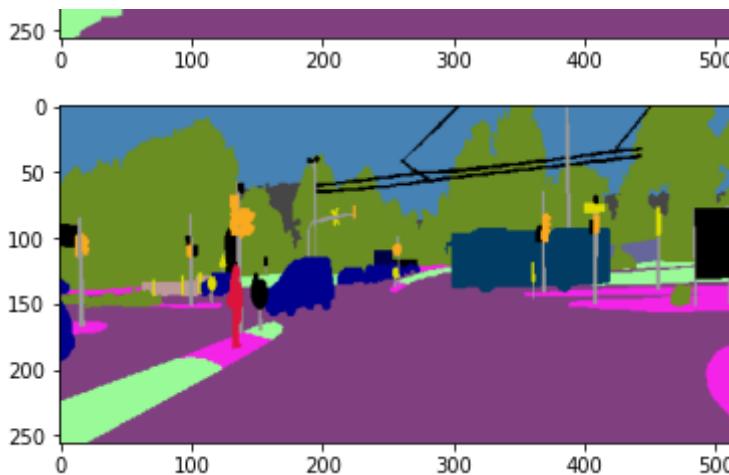
Test set Number: 27

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.2424309001515683

Dice: 0.3178963208223133



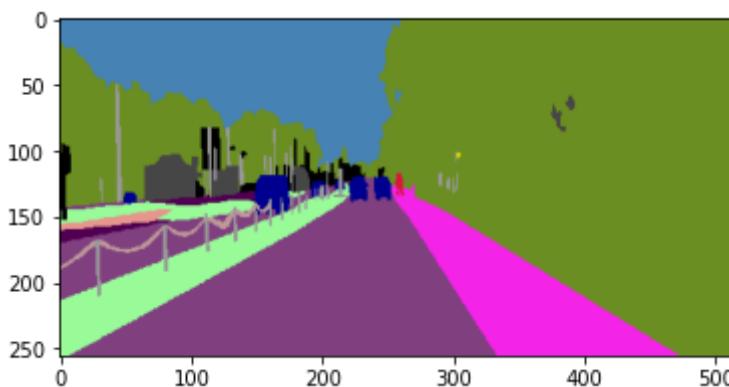
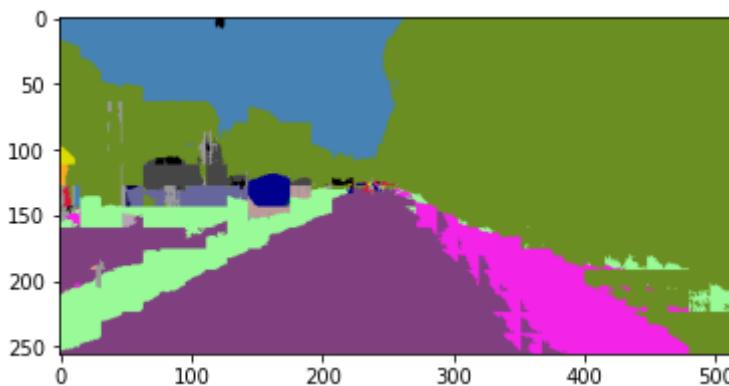
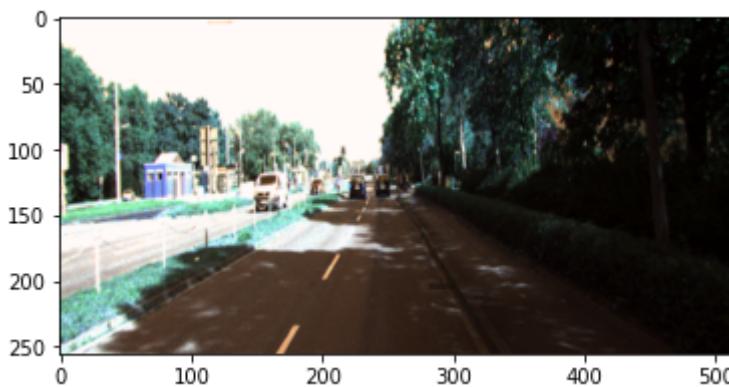


Test set Number: 28

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.2476639916627642

Dice: 0.2975545443083053



Test set Number: 29

Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

IoU: 0.3577634232905005

Dice: 0.44285500528200034

